

Izmir Institute of Technology

CENG 461 – Artificial Intelligence

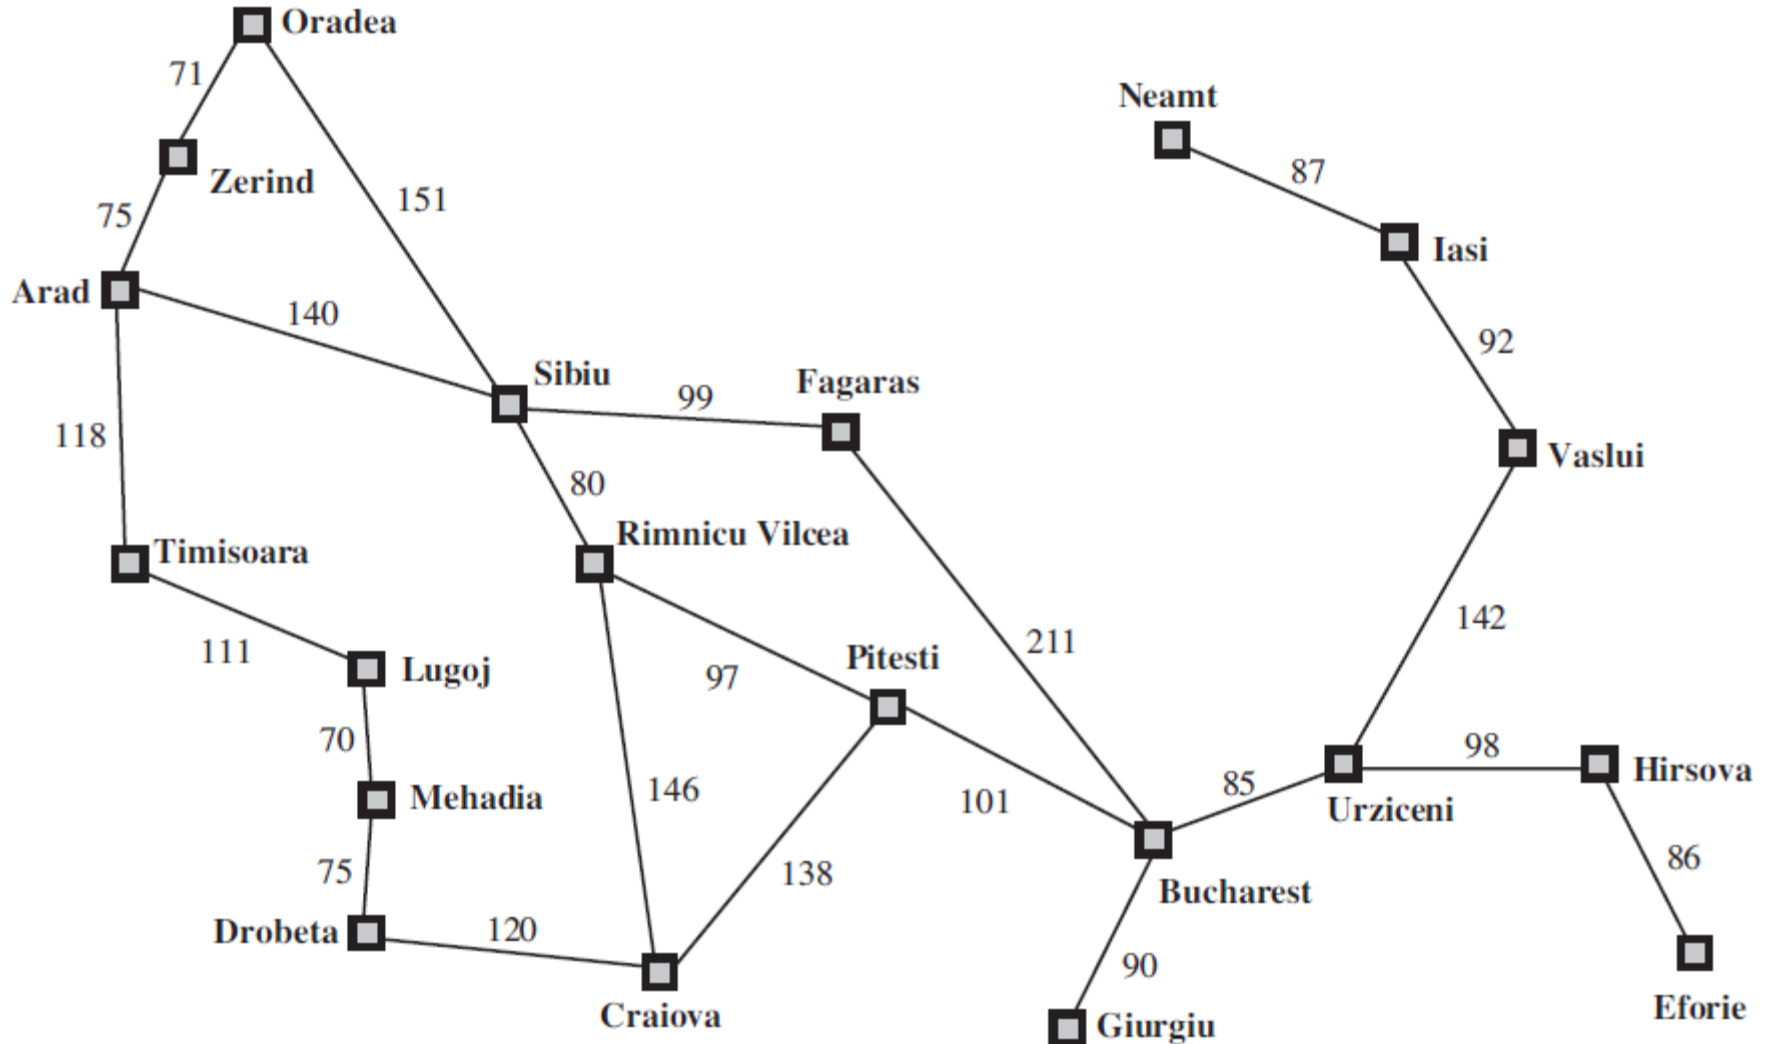
Solving Problems by Searching

Introduction

- ▶ We will consider problems that can not be solved by a single action.
- ▶ The solution will require a sequence of actions.
- ▶ At each step we will have multiple possible actions with possibly different but known costs.
- ▶ The difficulty lies in the large set of possible states that we need to consider to find the optimal sequence of actions to reach our goal.



Route Finding



8 Puzzle

7	2	4
5		6
8	3	1

	1	2
3	4	5
6	7	8




Problem Definition

- ▶ The initial state
 - ▶ s_0
- ▶ Actions for each state
 - ▶ $\text{Actions}(s) \rightarrow \{ a_1, a_2, a_3, \dots, a_N \}$
- ▶ Result state for each (state, action) pair
 - ▶ $\text{Result}(s, a) \rightarrow s'$
- ▶ Goal predicate
 - ▶ Checks whether we have reached a state or not, $\text{GoalTest}(s) \rightarrow \text{T|F}$
- ▶ Path cost for a sequence of actions from a starting state to an end state
 - ▶ If the path cost is the sum of step costs along the path, then we need to know the step cost for each (state, action) pair: $\text{StepCost}(s, a, s')$
- ▶ Solution/Optimal Solution



Vacuum-world



8-Puzzle

7	2	4
5		6
8	3	1

Initial State

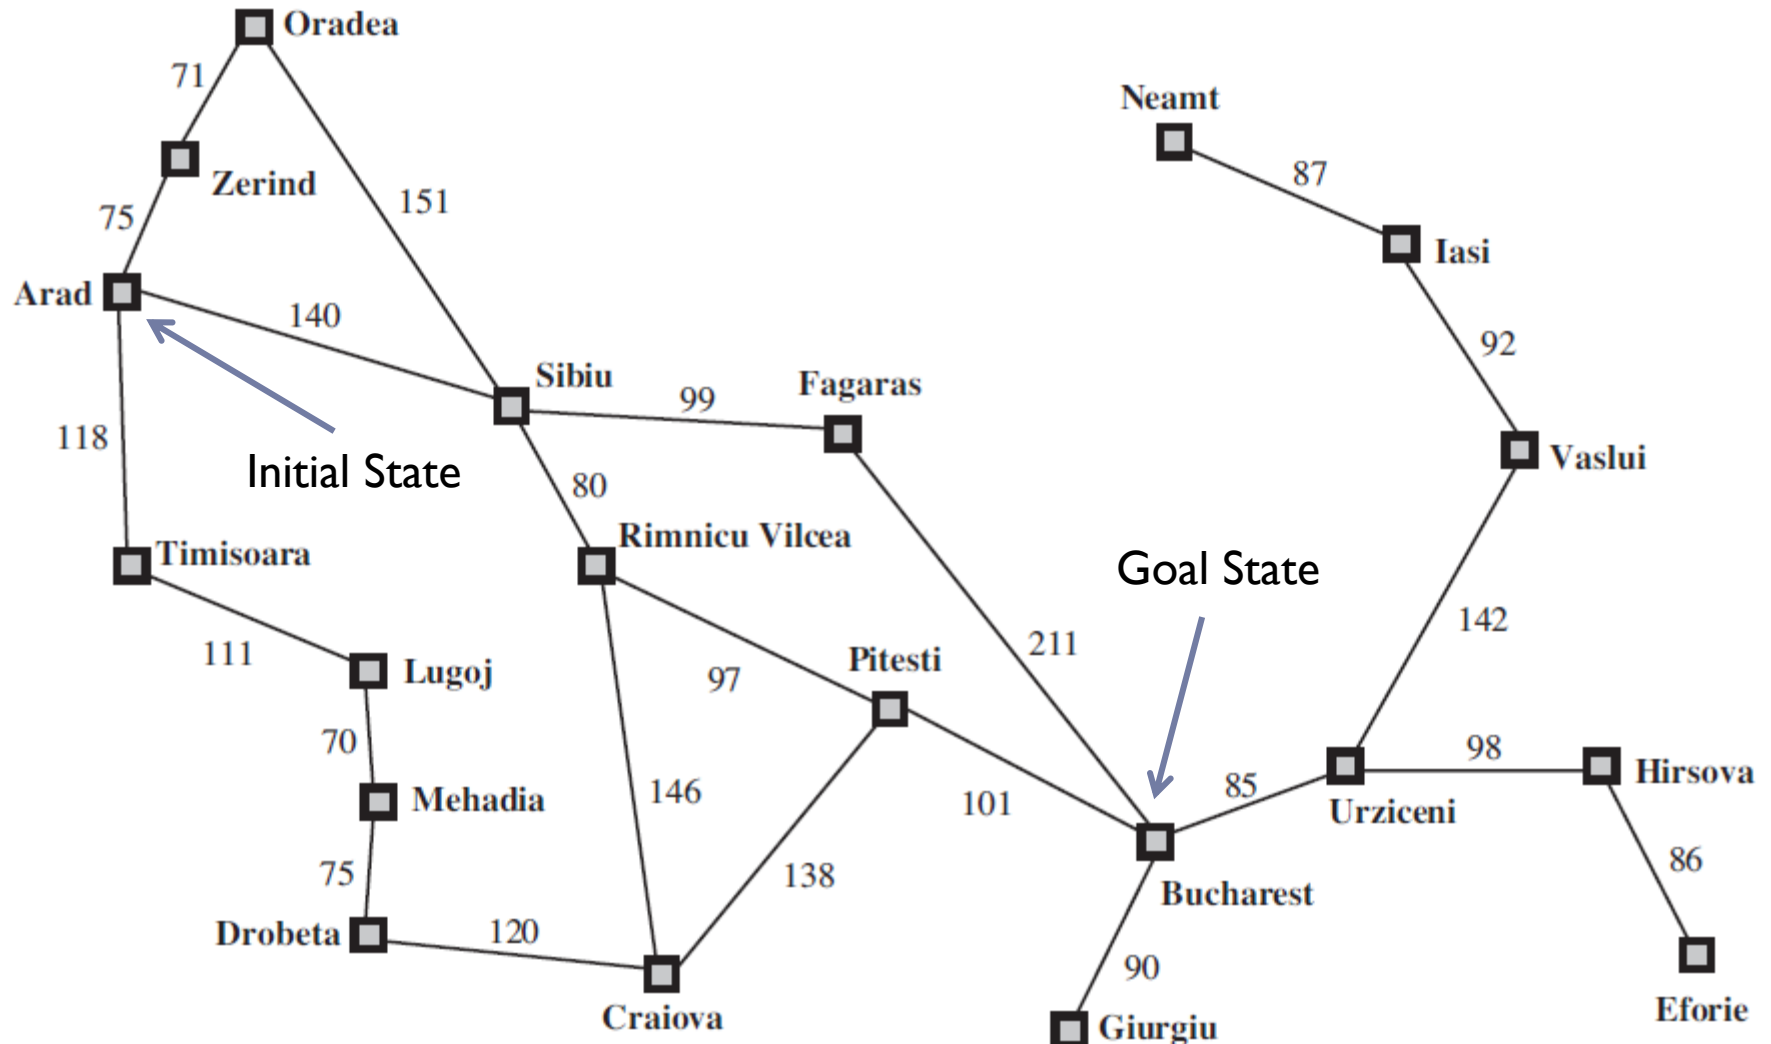
	1	2
3	4	5
6	7	8

Goal State

States, Actions, Path cost, Optimal solution ?



Route Finding



- States, Actions, Path cost, Optimal solution ?

Searching

- ▶ A search tree is used for the search.
 - ▶ The root node is the initial state.
 - ▶ Each branch corresponds to an action taken at the parent node.
 - ▶ For each node, we need to check if we are in a goal state, and if not expand the tree by applying all possible actions for that node leading to child nodes.
 - ▶ A leaf node has no children. The set of all leaf nodes available for expansion at any given point is called the **frontier**.
- ▶ For a search problem we have
 - ▶ Frontier States
 - ▶ Explored States
 - ▶ Unexplored States



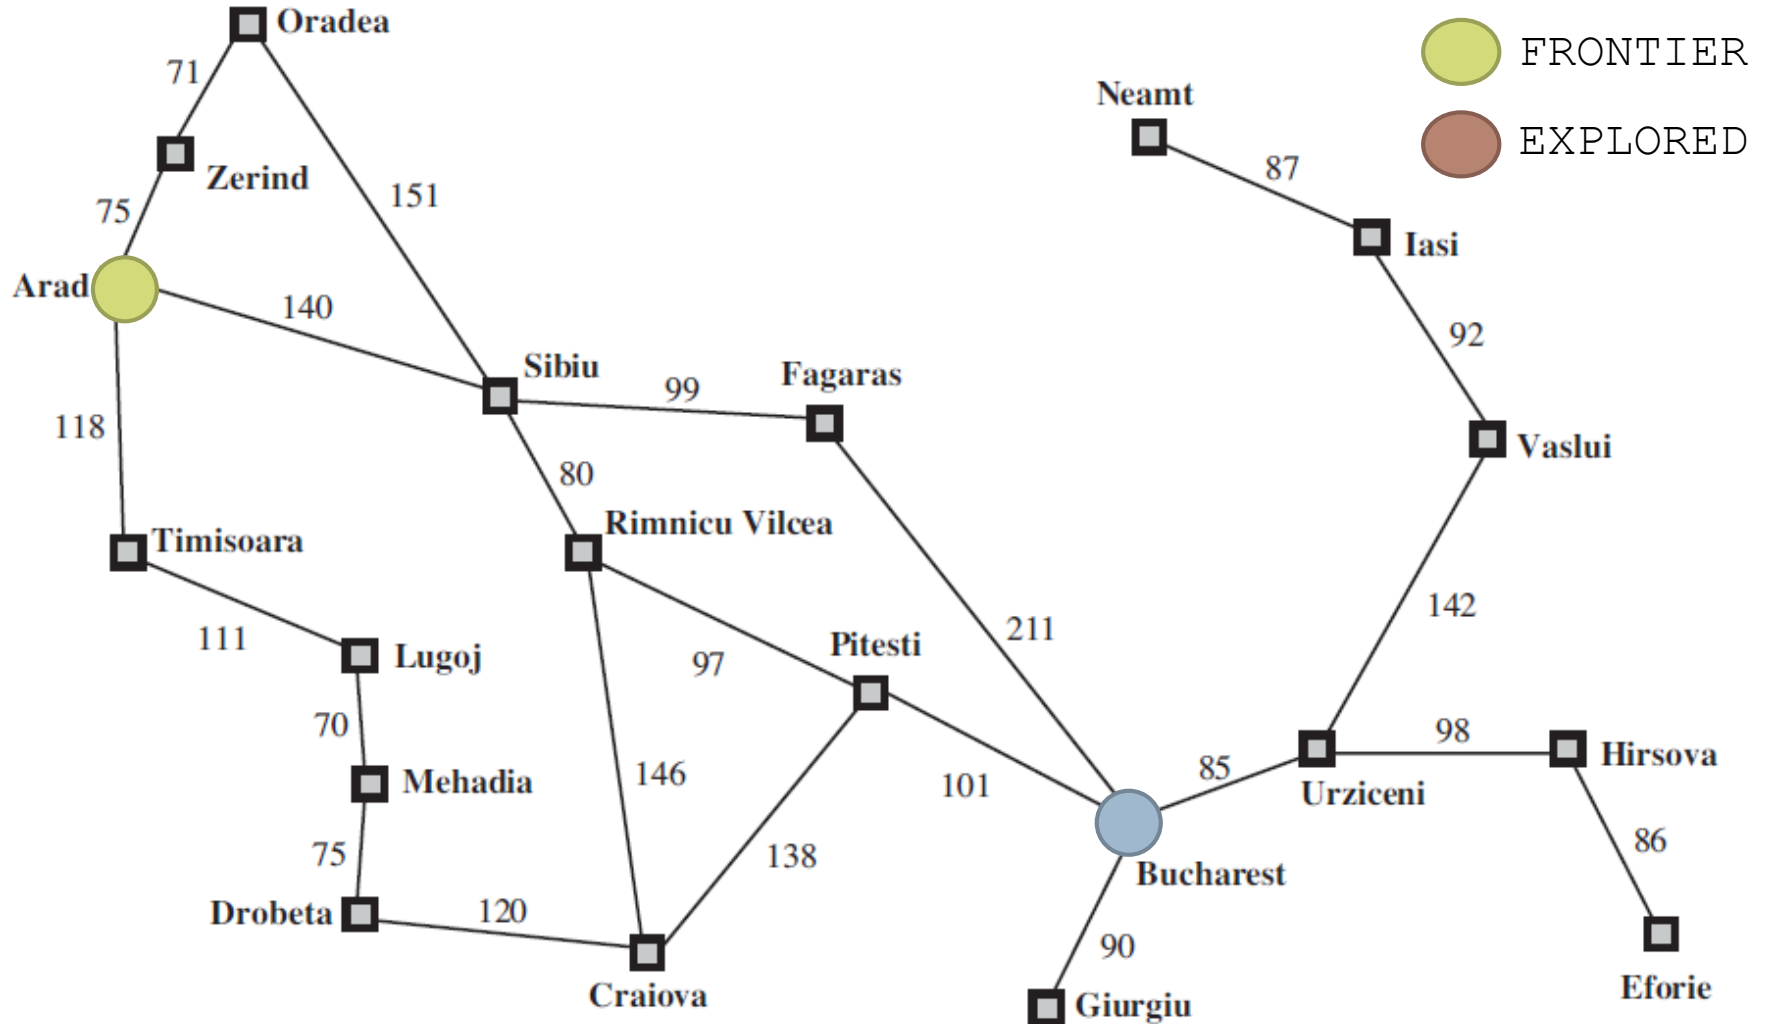
Tree Search

- ▶ At each step we need to choose which node to expand first (indicated by the blue line in the pseudo-code below).
- ▶ There are multiple ways to decide and this choice leads to different strategies.

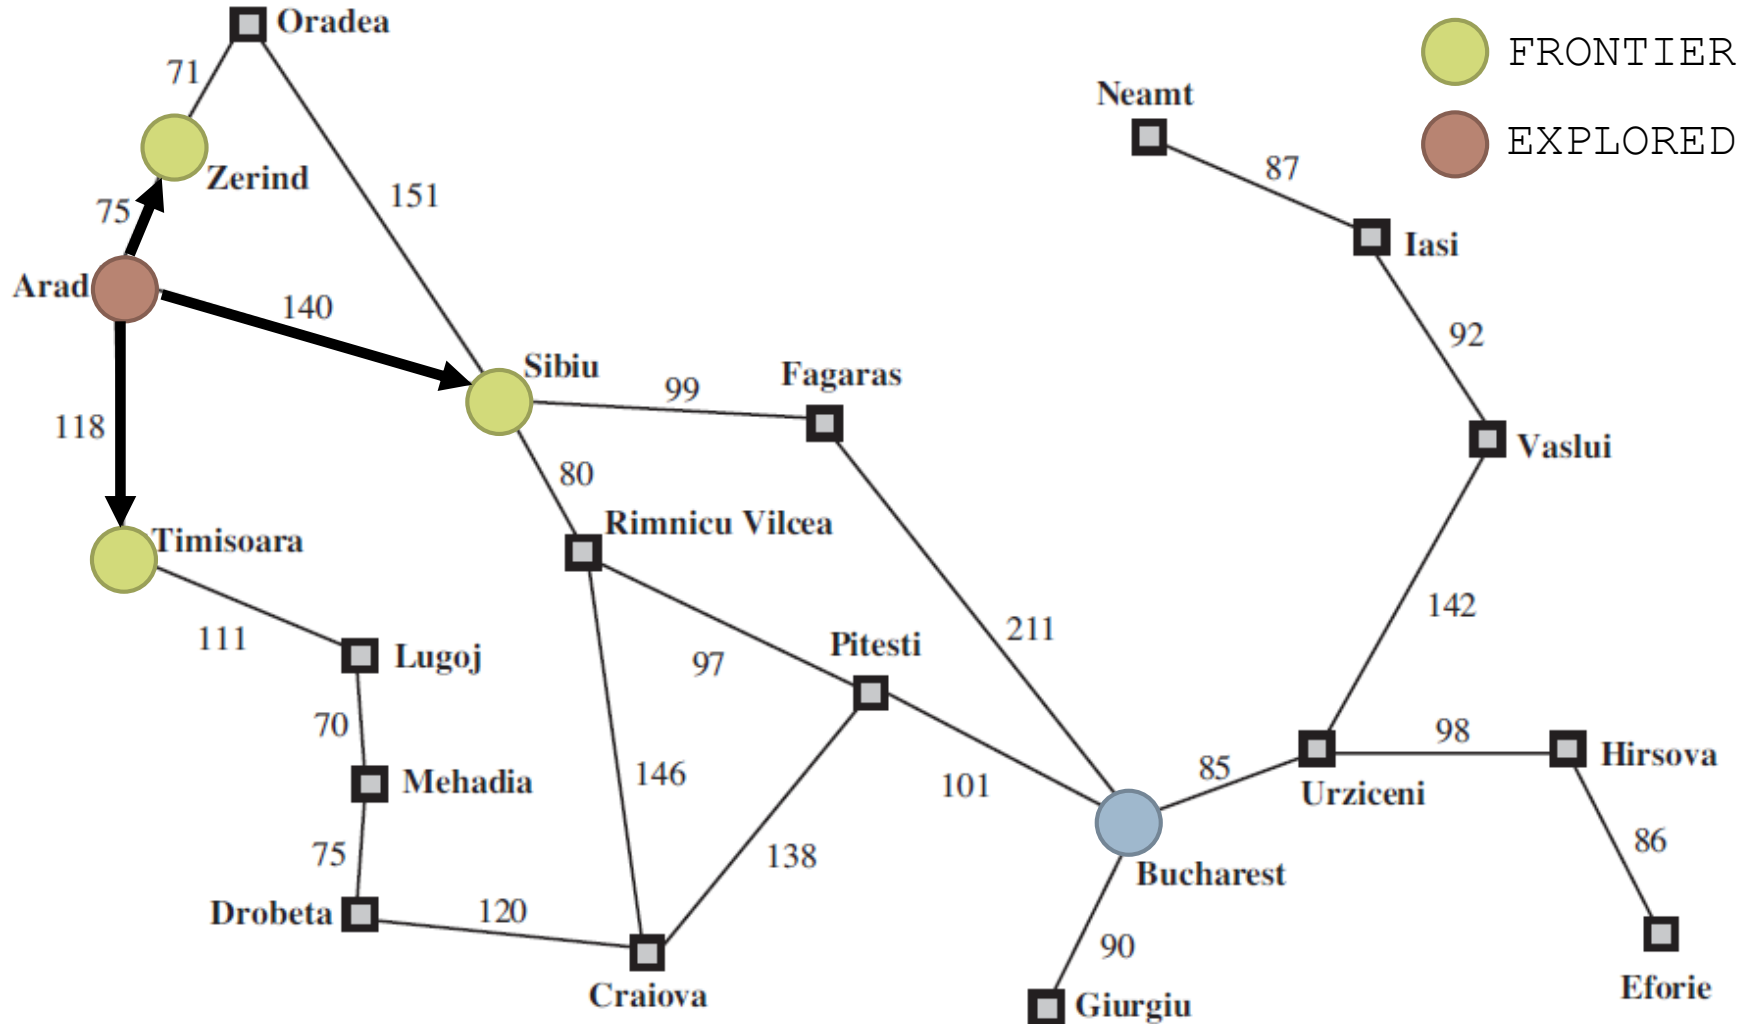
```
function TREE-SEARCH(problem) %returns a solution, or failure
  frontier = initial state %initialize the frontier
  loop do
    if frontier is empty then
      return failure
    choose a leaf node (s) and remove it from frontier
    if node s contains a goal state then
      return the corresponding solution
    expand node s by applying possible actions (result(s,a))
    add the resulting nodes to the frontier
```



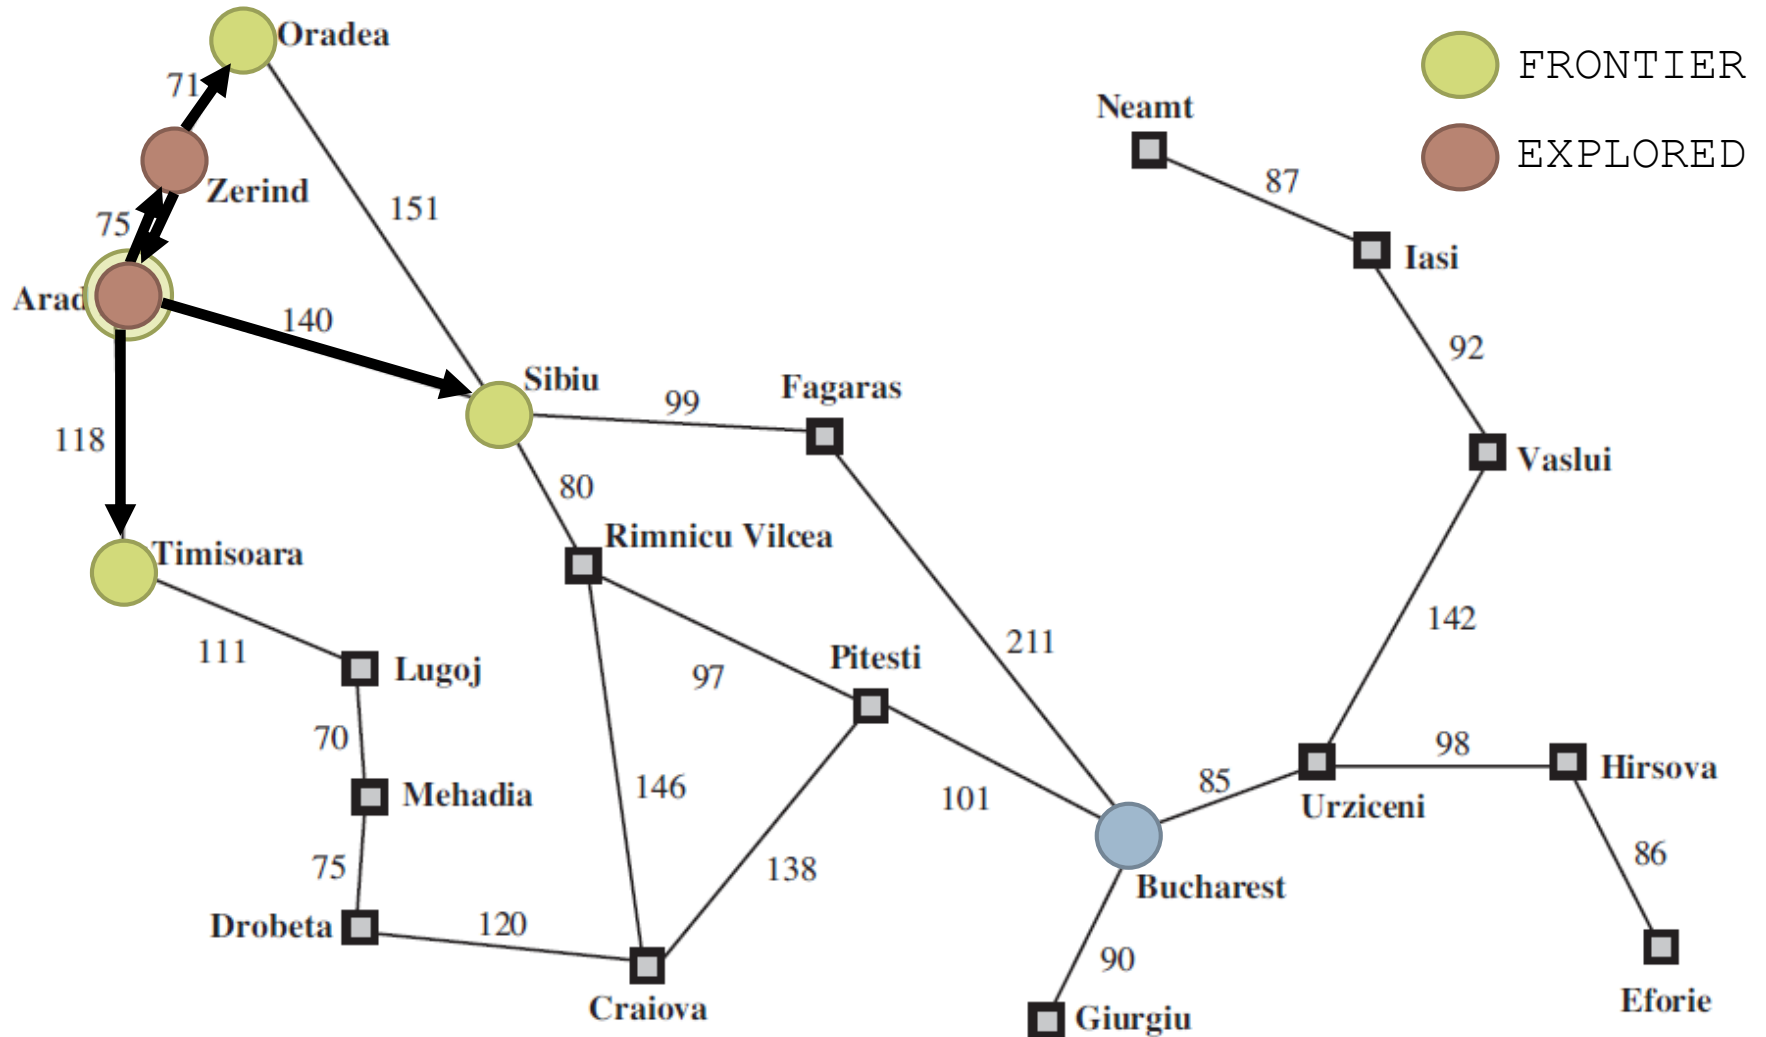
Tree Search for Route Finding




Tree Search for Route Finding




Tree Search for Route Finding



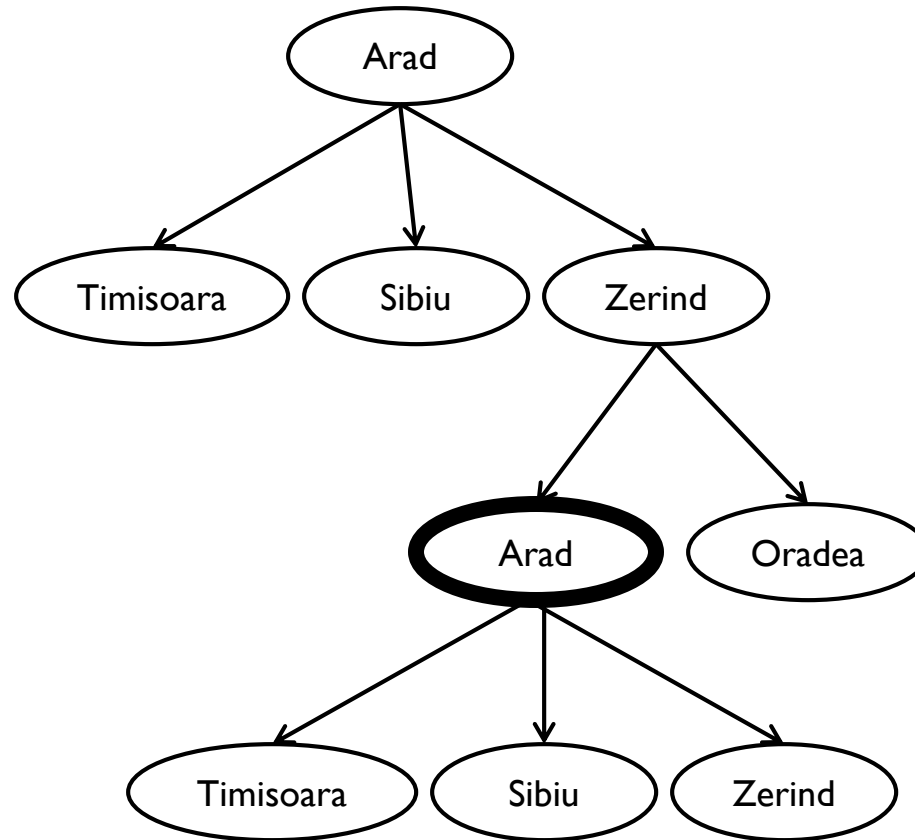
Tree Search for Route Finding



Tree Search for Route Finding



Tree Search for Route Finding



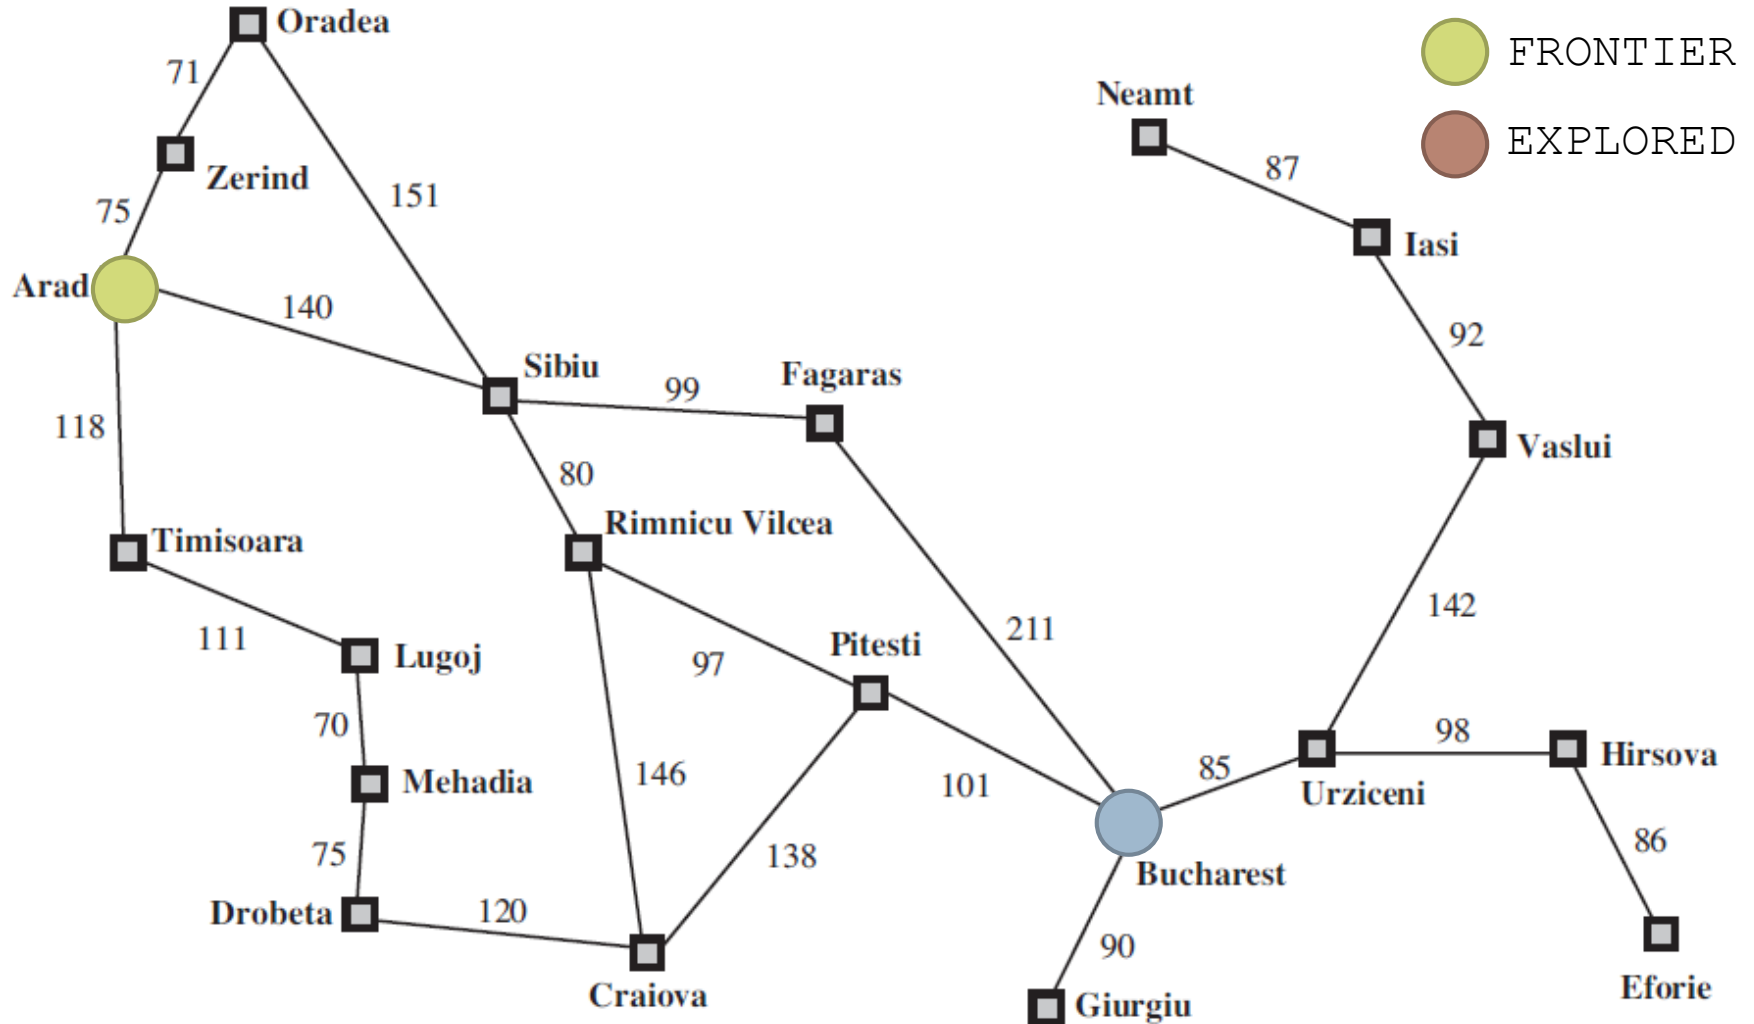
Avoiding Loops with Graph Search

- ▶ When it is possible to get back to a previously visited states, the tree-search algorithm explores it again.
- ▶ Such loops can be avoided by keeping a history of explored states.

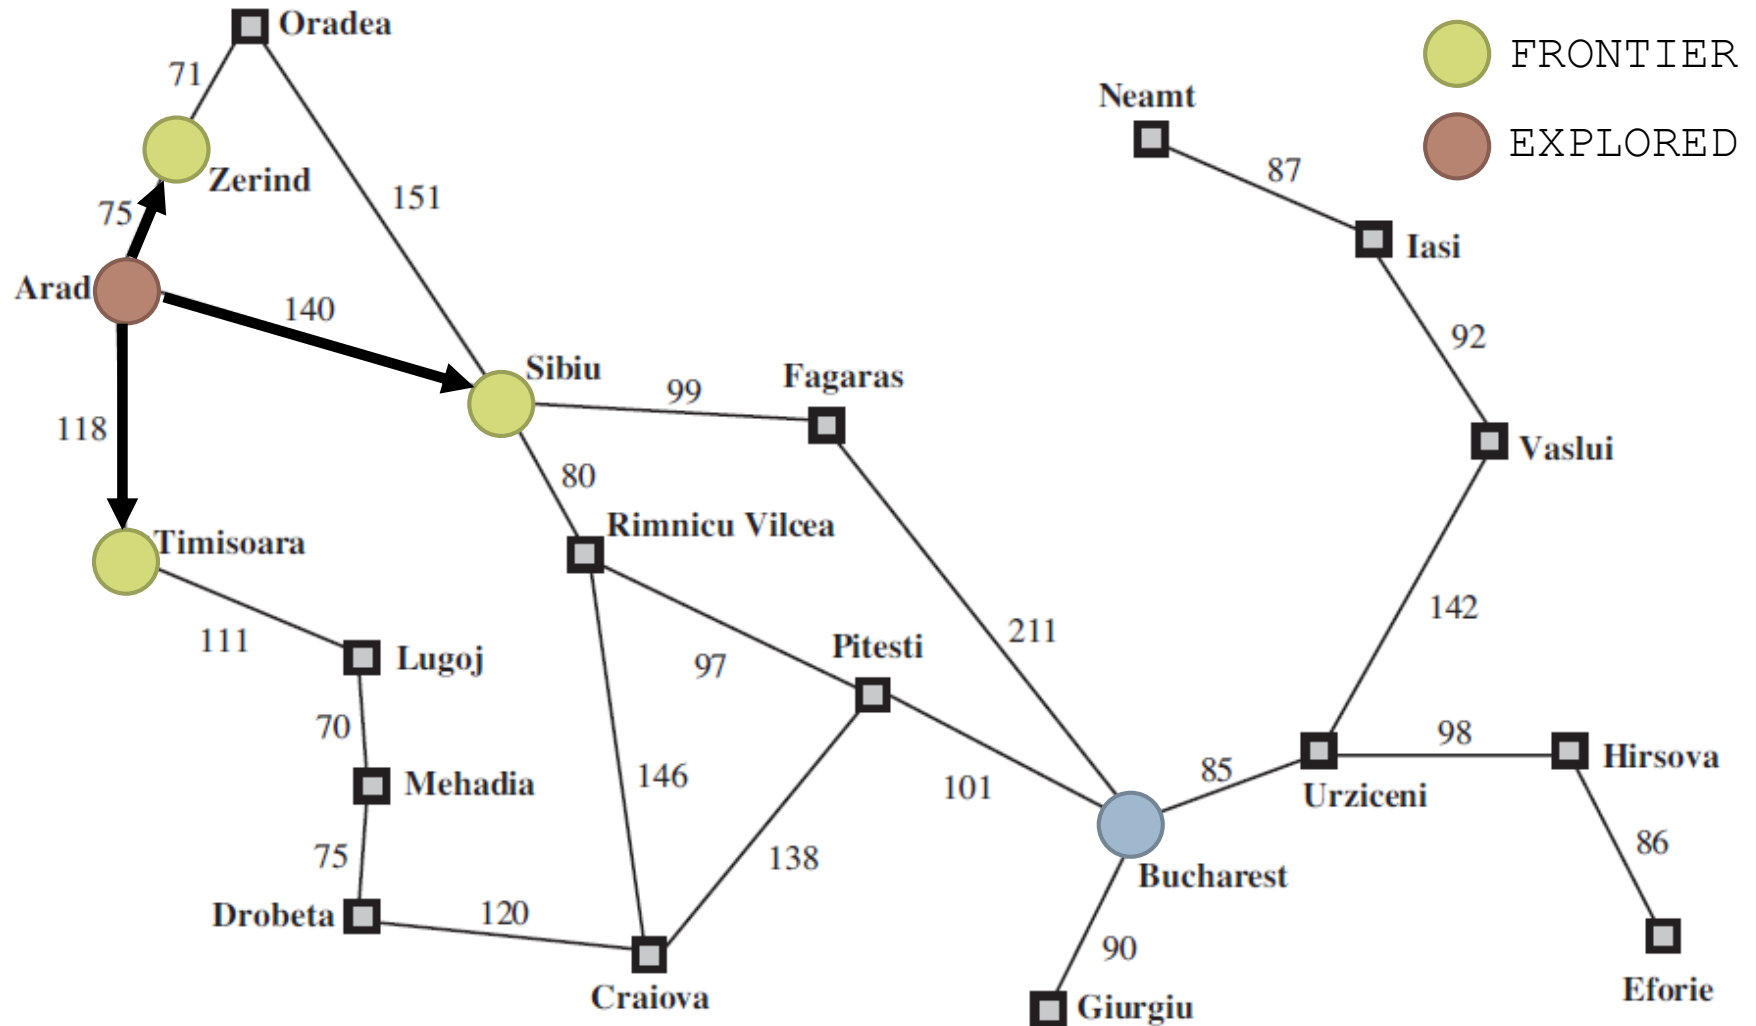
```
function GRAPH-SEARCH(problem) %returns a solution, or failure
    frontier = initial state %initialize the frontier
    explored = {} %initialize the explored set to be empty
    loop do
        if the frontier is empty then
            return failure
        choose a leaf node (s) and remove it from frontier
        add node s to the explored
        if node s contains a goal state then
            return the corresponding solution
        expand node s by applying all possible actions (result(s,a))
        if the resulting node is not in frontier or explored then
            add the resulting nodes to frontier
```

▶ * Blue lines indiate the difference between graph-search and tree-search

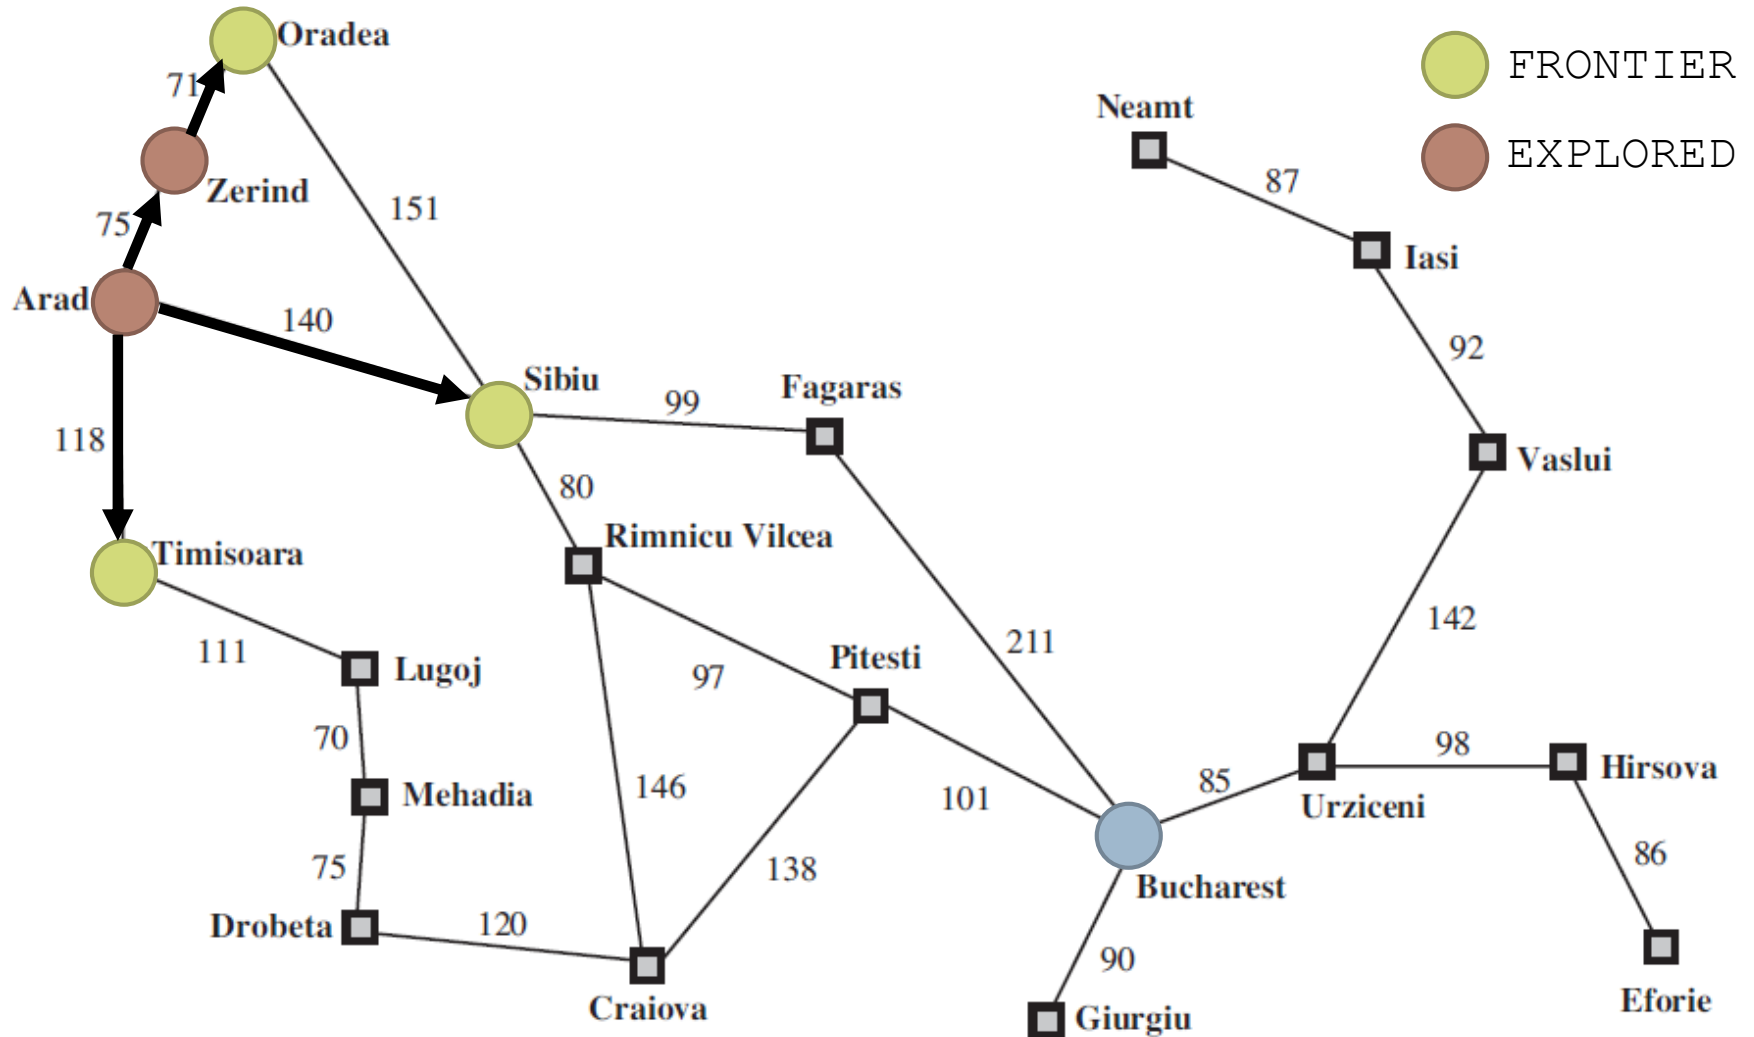
Graph Search for Route Finding



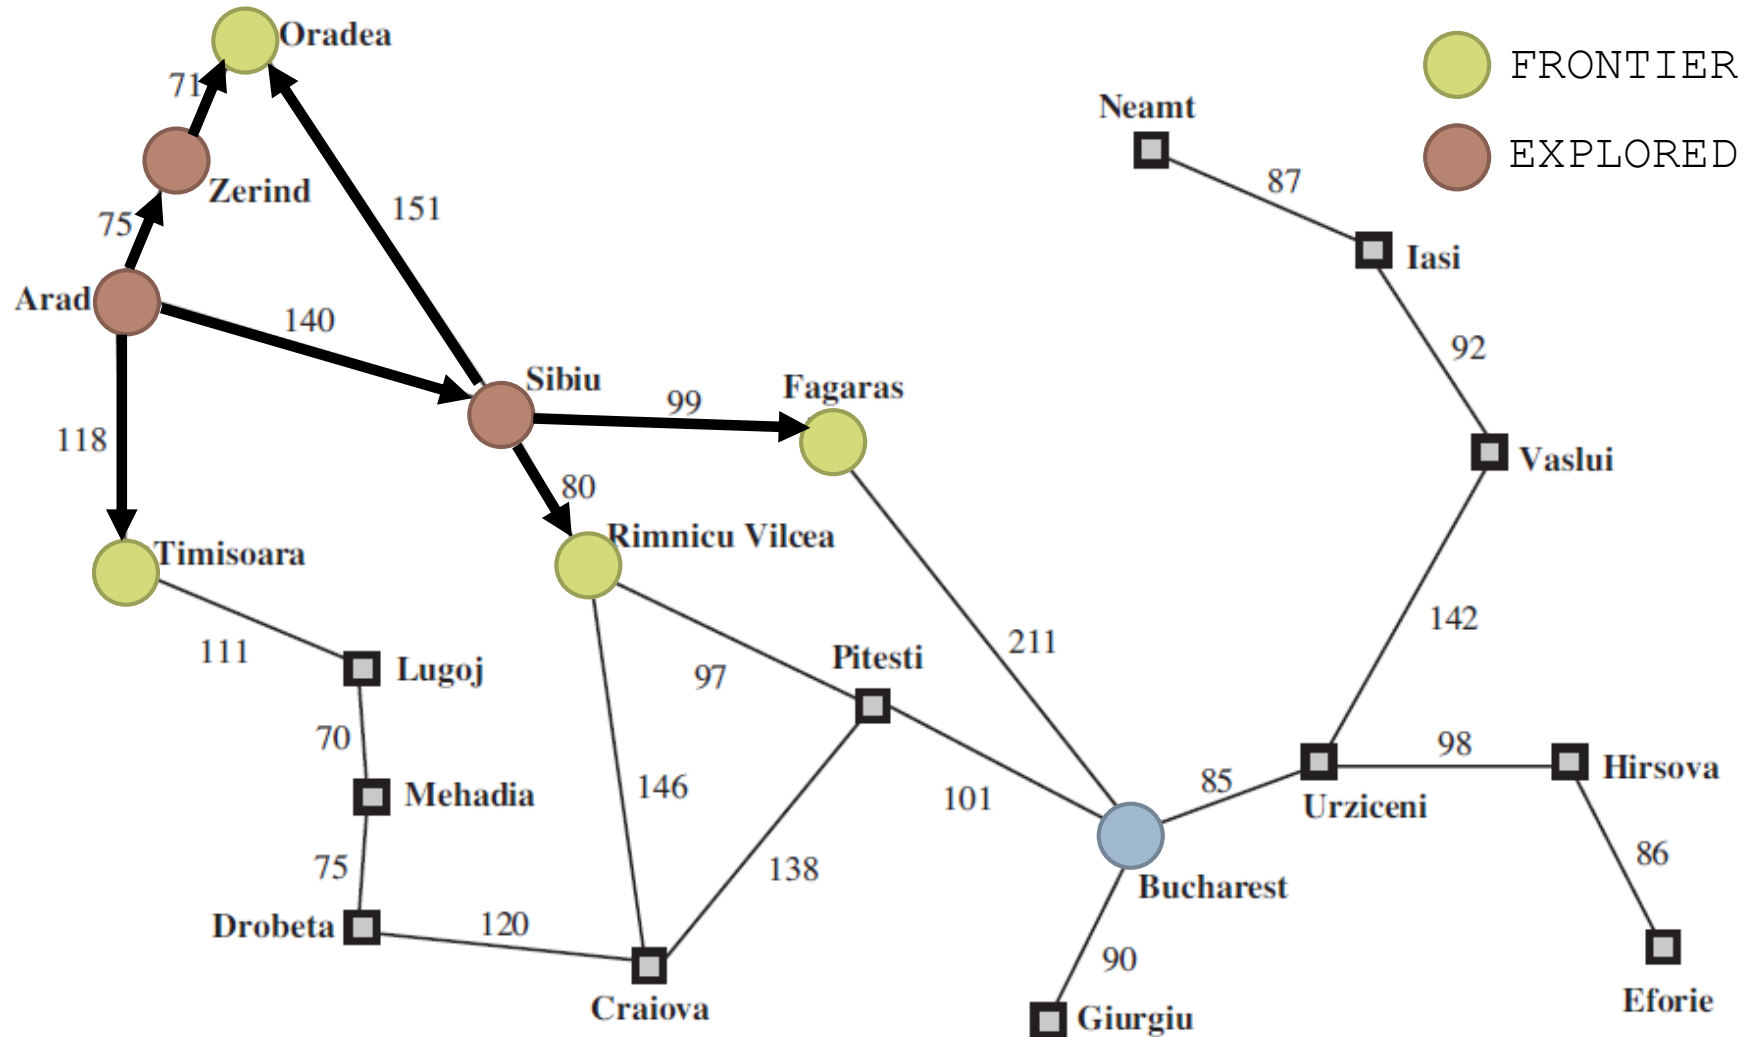
Graph Search for Route Finding



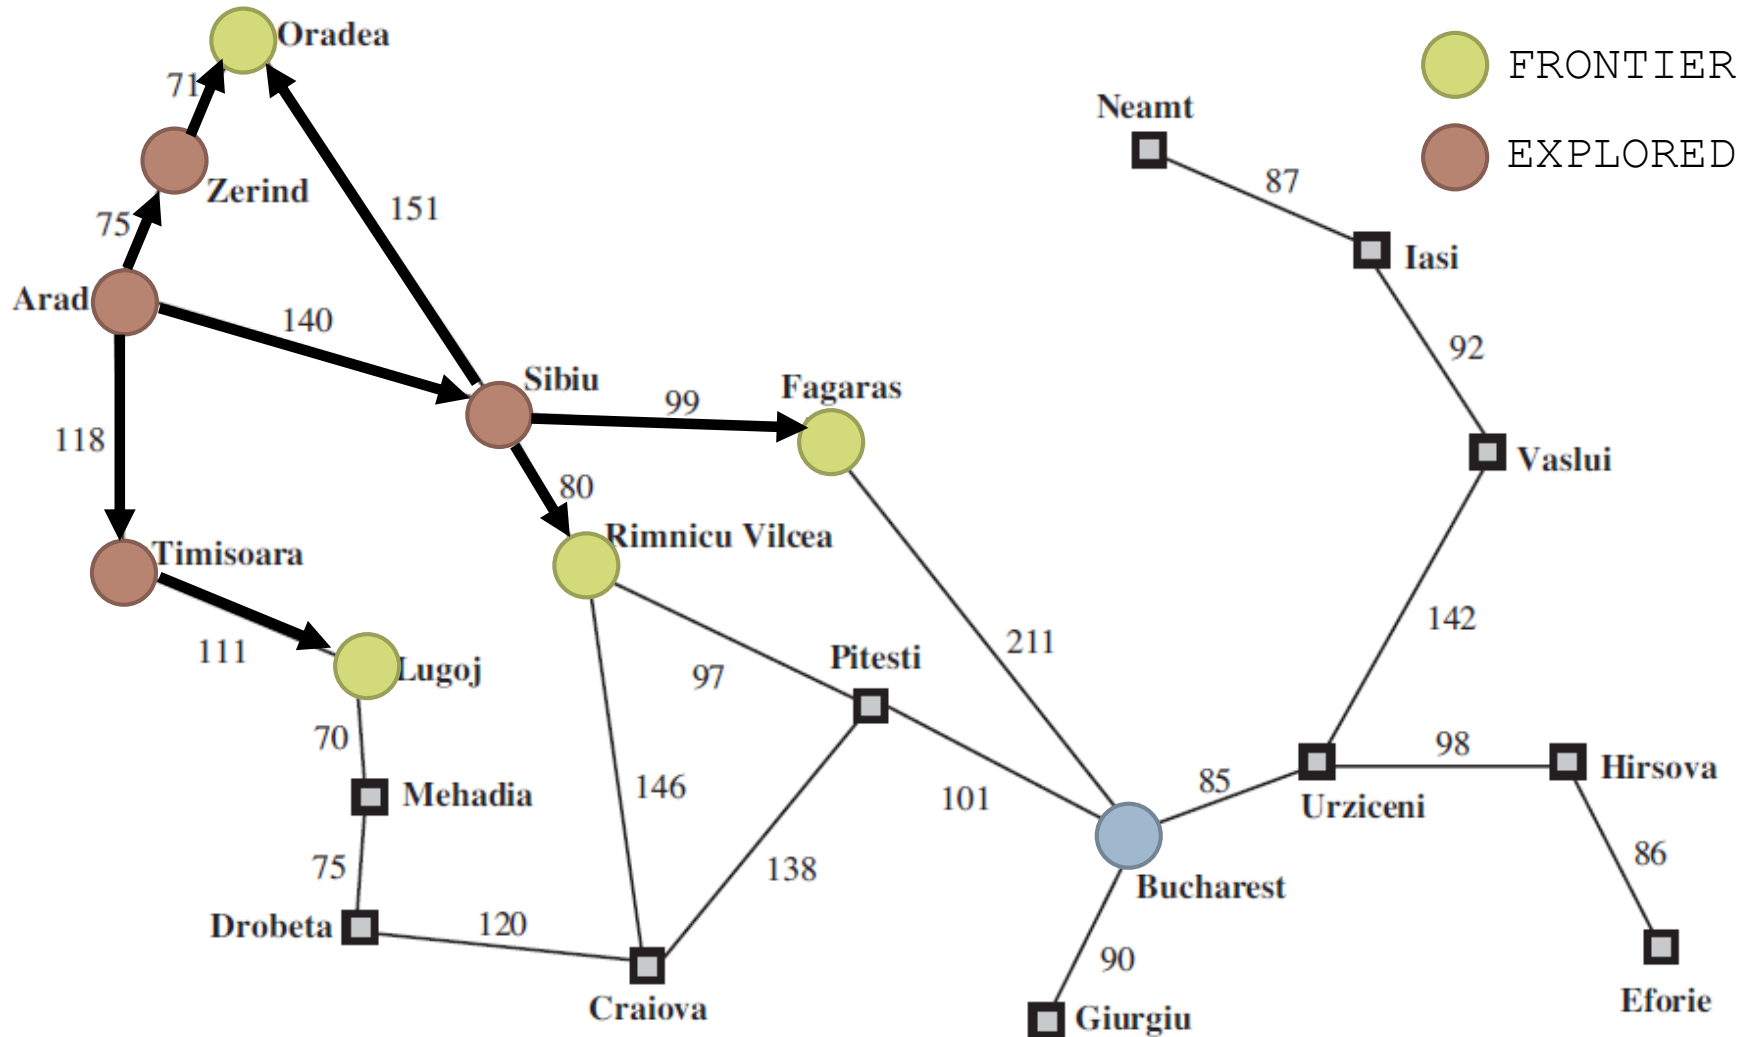
Graph Search for Route Finding




Graph Search for Route Finding



Graph Search for Route Finding




Graph Search for Route Finding



Implementing Search Algorithms

- ▶ For each node, we need
 - ▶ **state**: the state that the node corresponds to
 - ▶ **parent**: the node that generated this one, we need this to backtrack from a goal state.
 - ▶ **action**: the action at the parent node.
 - ▶ **path_cost**: the cost from initial state to this node.



Implementing Search Algorithms

- ▶ **Make sure to distinguish between states and nodes**

Nodes lie on a path whereas a state corresponds to a configuration of the world. Same state can occur at different nodes. E.g. You can reach a city using different paths.

- ▶ **The frontier is usually stored in a queue** (the algorithm can easily choose the next node to expand according to its preferred strategy):
 - ▶ It could be a regular FIFO queue.
 - ▶ It could be a stack.
 - ▶ It could be a priority queue.
- ▶ **We also need to store the explored set**



Measuring Performance

- ▶ **Completeness**

- ▶ Is it guaranteed to find a solution?

- ▶ **Optimality**

- ▶ Is it guaranteed to find the optimal solution?

- ▶ **Time Complexity**

- ▶ How does the running time and problem size related?

- ▶ **Space Complexity**

- ▶ How does the required memory and problem size related?



Uninformed Search Strategies

- ▶ When we do not know anything about relative suitability of expanding non-goal states, we perform an uninformed (blind) search.
- ▶ The strategies below are distinguished by the order in which nodes are expanded.
 - ▶ Breadth-First Search
 - ▶ Uniform-Cost Search
 - ▶ Depth-First Search
 - ▶ Depth-Limited Search
 - ▶ Iterative Deepening Depth-First Search
 - ▶ Bidirectional Search



Breadth-First Search

- ▶ Expand the root node first, then its successors, then all the successors of the root node are expanded next, then their successors, and so on.
- ▶ All the nodes are expanded at a given depth (n) in the search tree before any nodes at the next level ($n+1$) are expanded.
- ▶ The frontier can be implemented as a FIFO queue.



Breadth-First Search (for reference)

function BREADTH-FIRST-SEARCH(problem) **returns** a solution, or failure

node \leftarrow a node with STATE = problem.INITIAL-STATE, PATH-COST = 0

if problem.GOAL-TEST(node.STATE) **then return** SOLUTION(node)

frontier \leftarrow a FIFO queue with node as the only element

explored \leftarrow an empty set

loop do

if EMPTY?(frontier) **then return** failure

 node \leftarrow POP(frontier) /* chooses the shallowest node in frontier */

 add node.STATE to explored

for each action **in** problem.ACTIONS(node.STATE) **do**

 child \leftarrow CHILD-NODE(problem, node, action)

if child.STATE is not in explored or frontier **then**

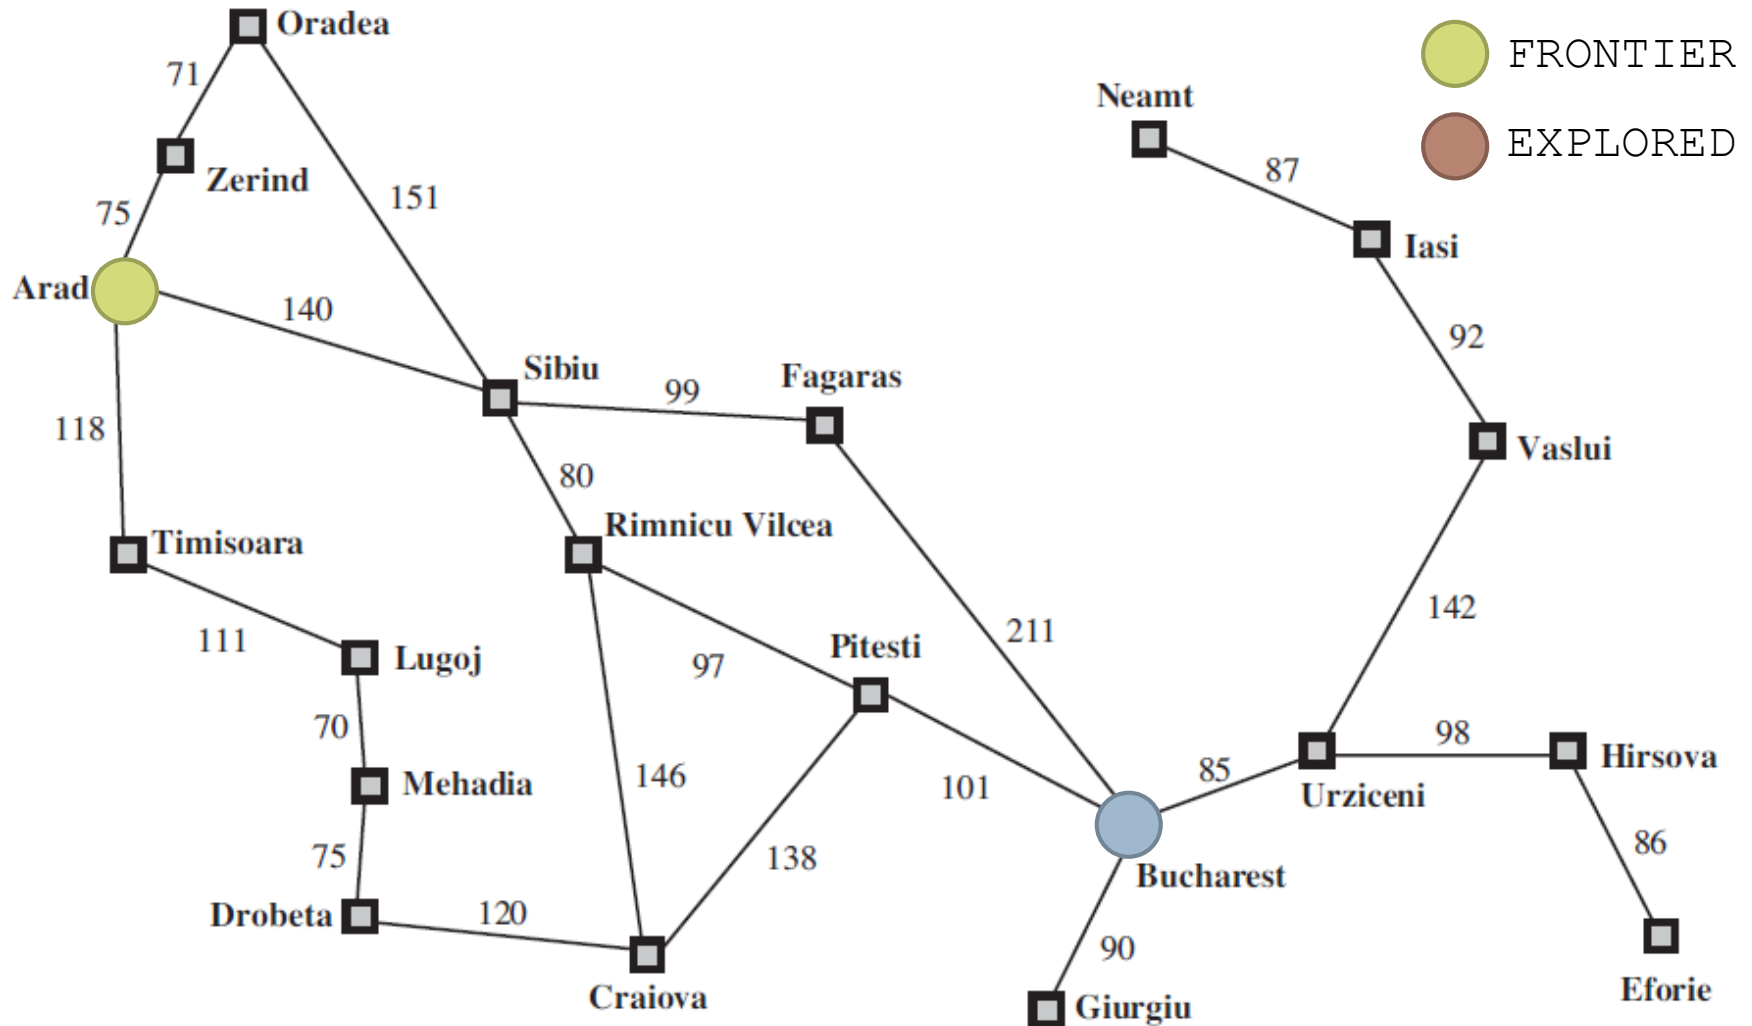
if problem.GOAL-TEST(child.STATE) **then**

return SOLUTION(child)

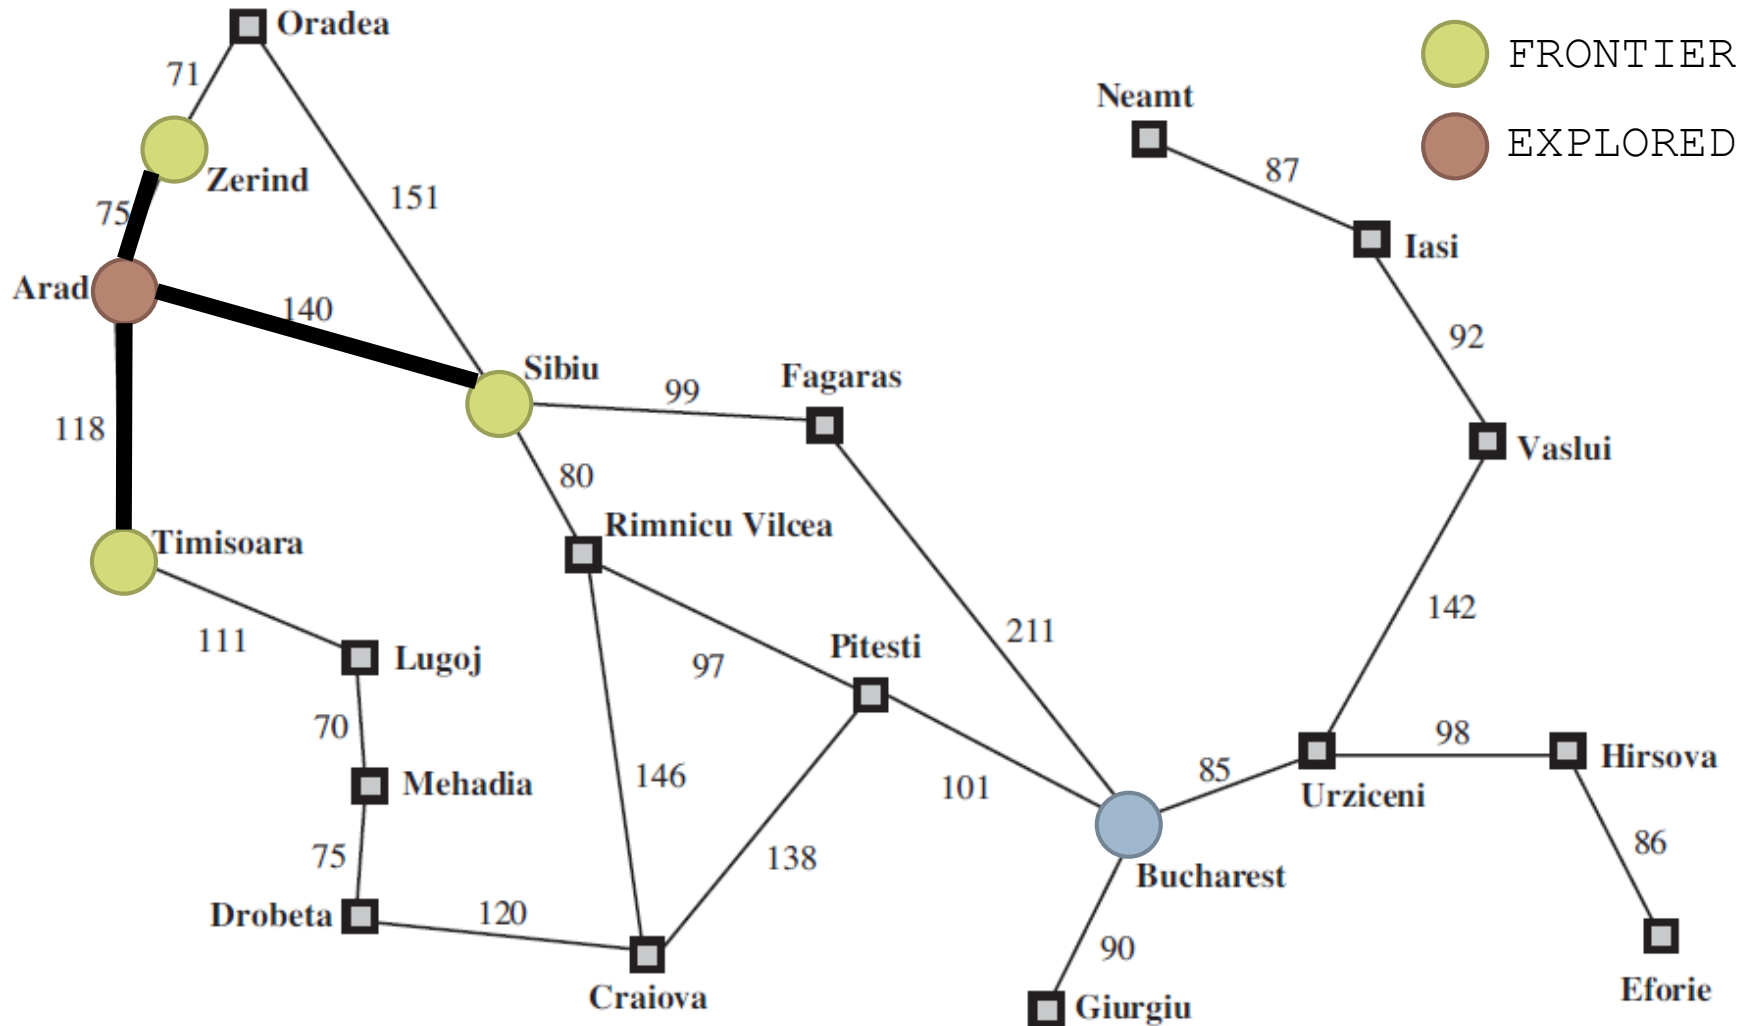
 frontier \leftarrow INSERT(child, frontier)



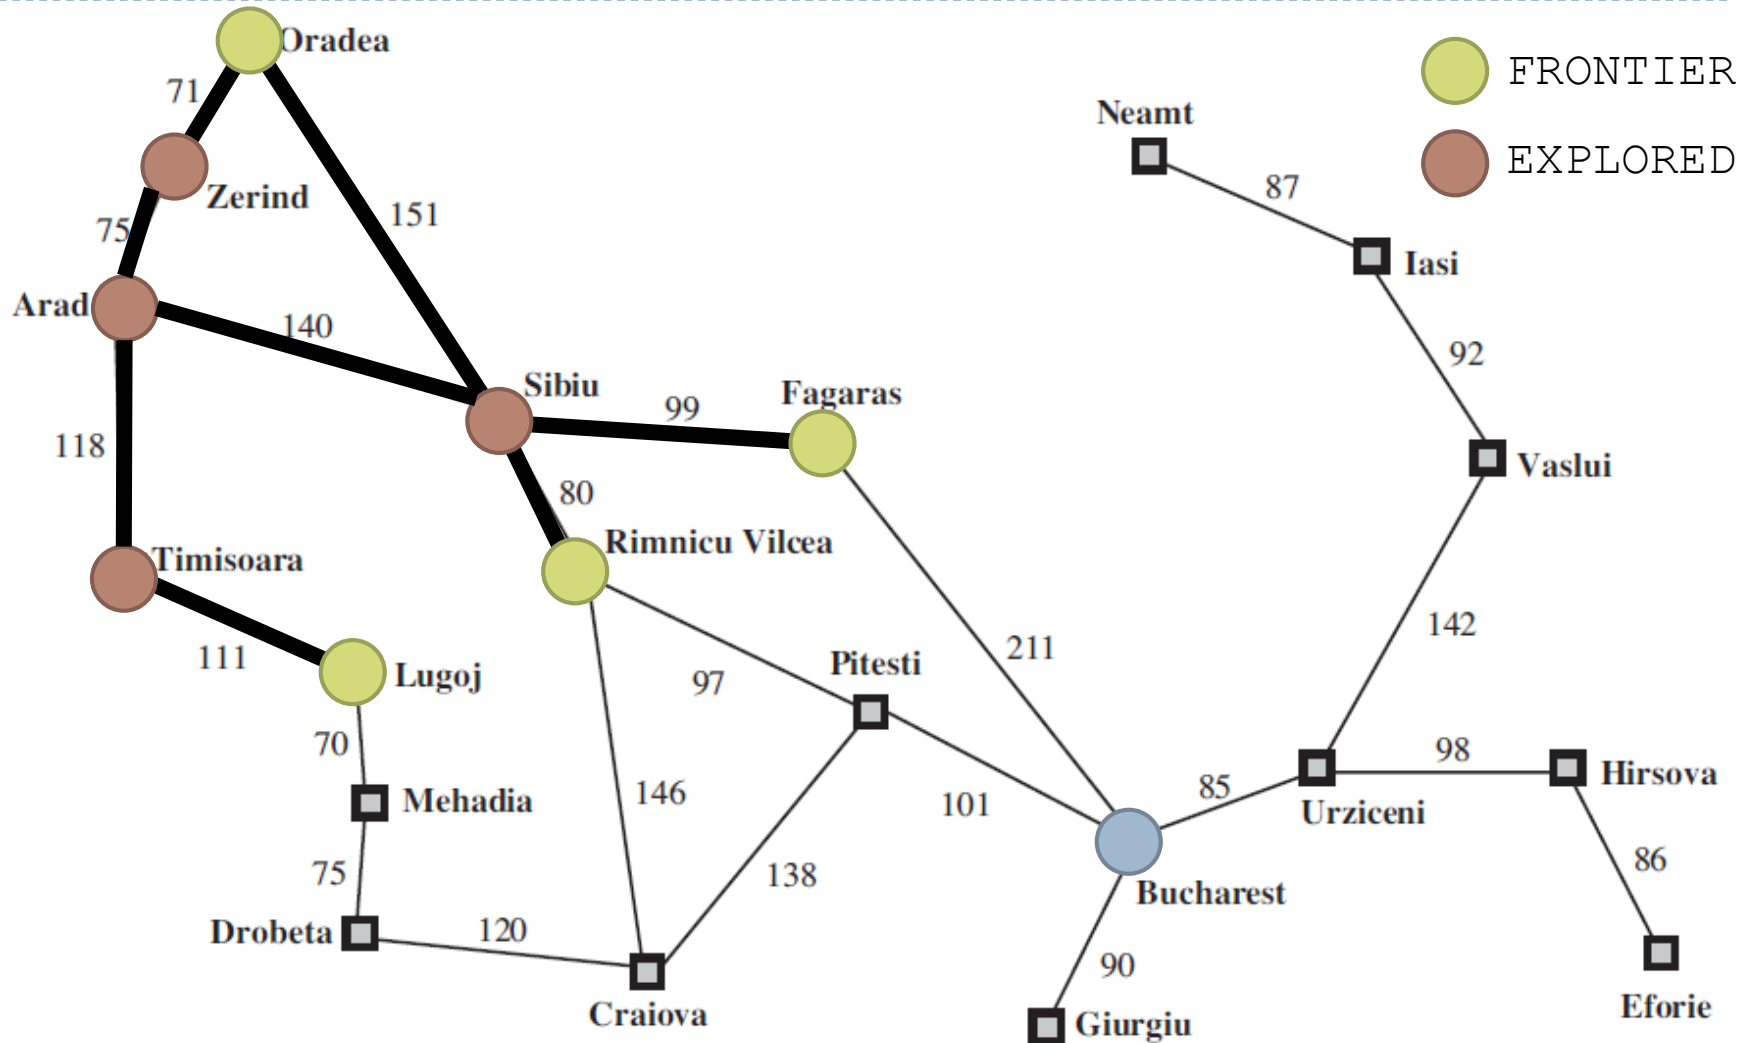
Breadth-First Search for Route Finding



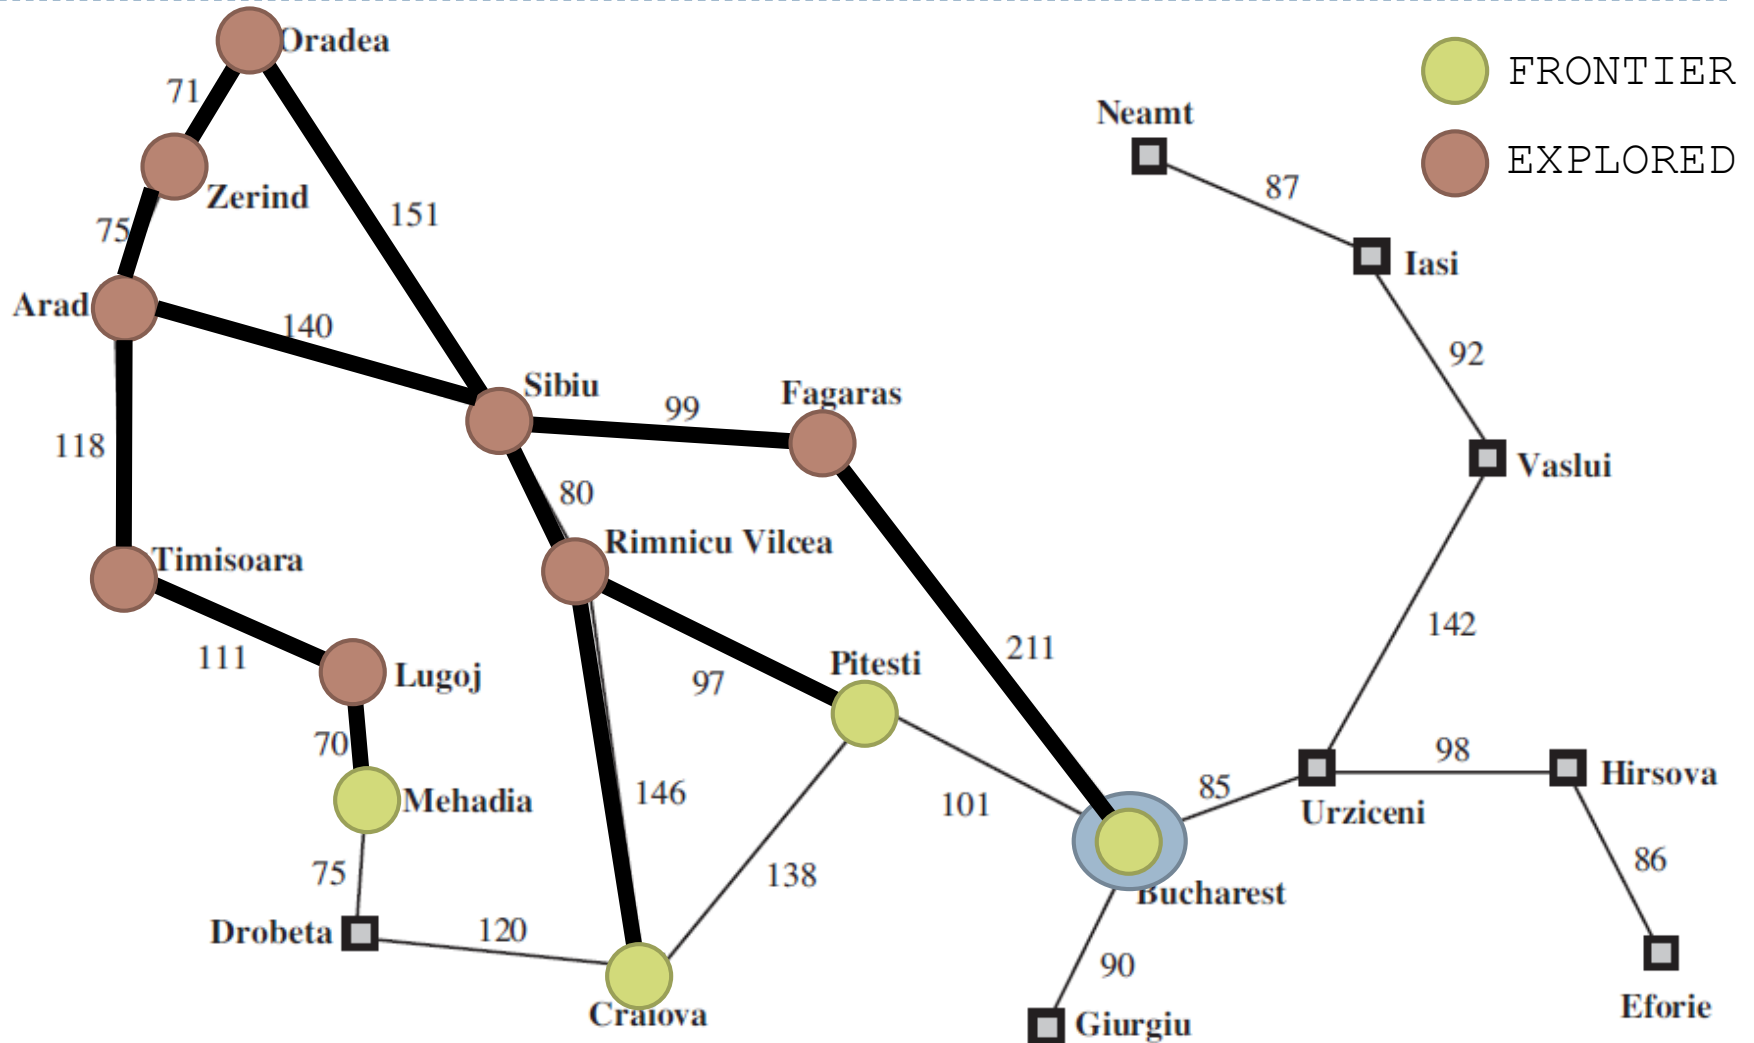
Breadth-First Search for Route Finding



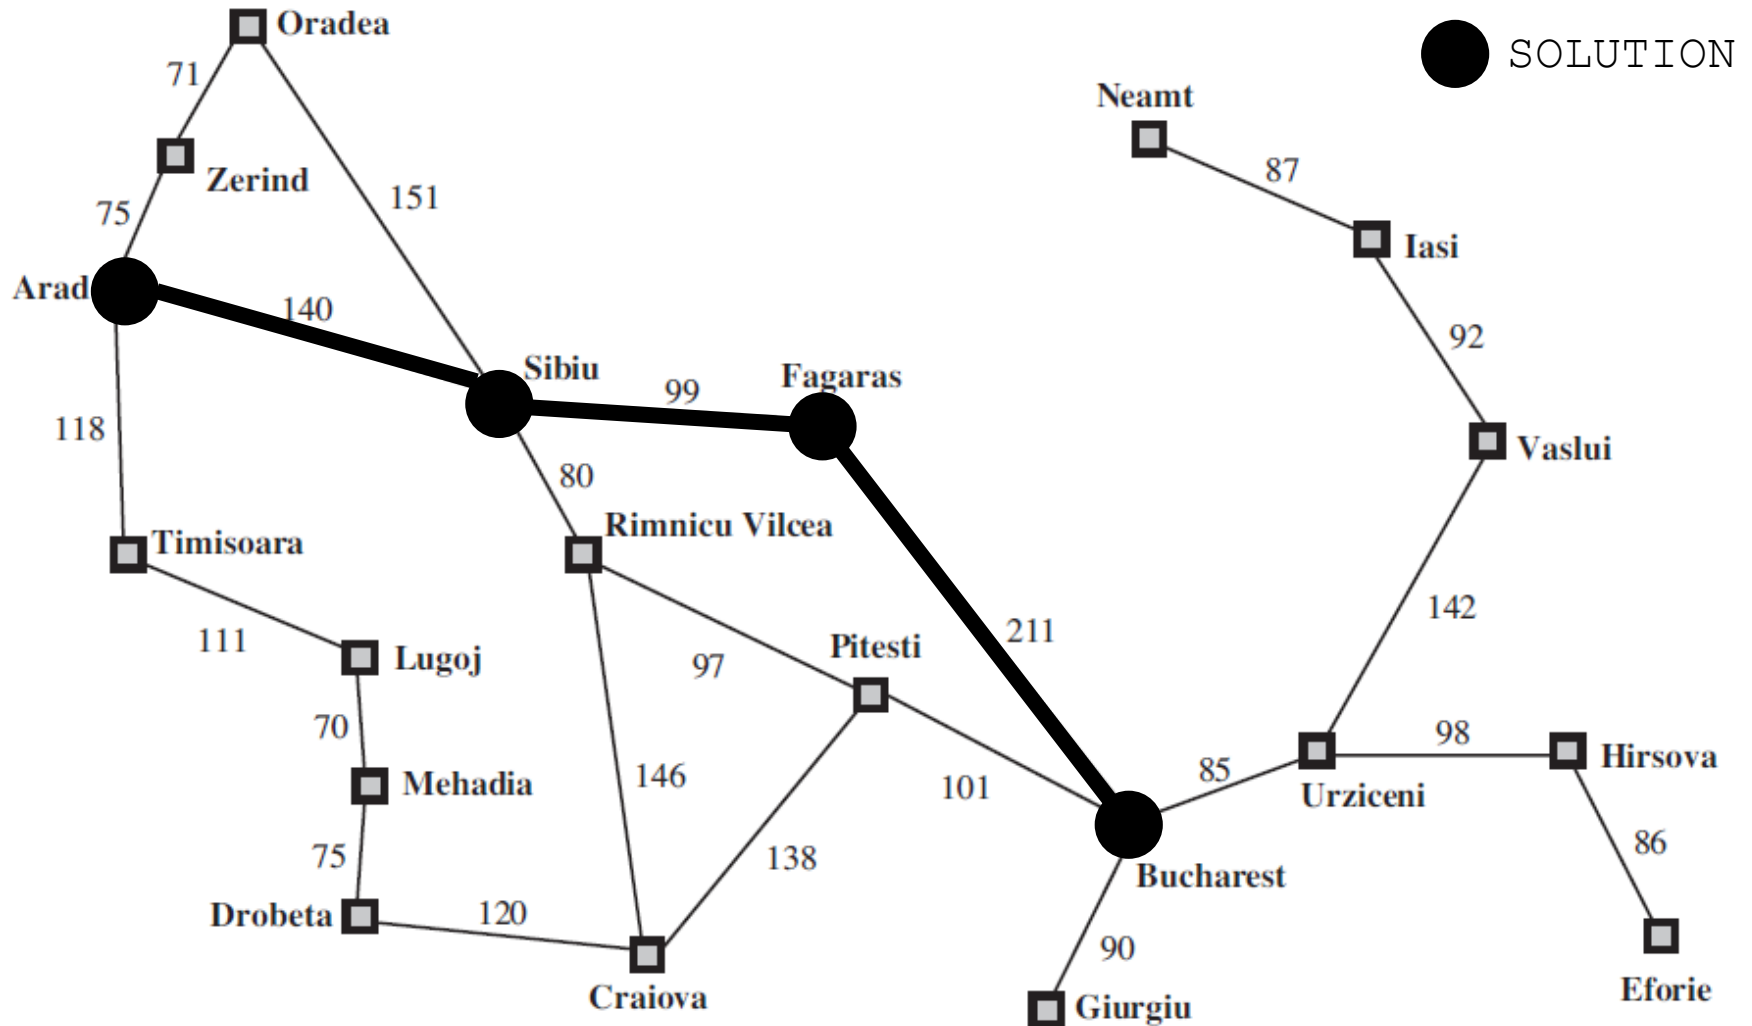
Breadth-First Search for Route Finding



Breadth-First Search for Route Finding



Breadth-First Search for Route Finding



► Is it the minimum cost path?


Breadth-First Search

- ▶ It is complete since if the goal state is at depth n , it is guaranteed to be explored after $n-1$.
- ▶ It is optimal only if the path cost is a non-decreasing function of the depth of the node, i.e. if a better solution can not exist deeper (e.g. all step costs are equal).

- ▶ Its time complexity is $O(b^d)$, where b is the branching factor and d is the solution depth.

If goal-state is checked after explored $\rightarrow O(b^{d+1})$

- ▶ Its space complexity is $O(b^d)$, i.e. the number of nodes in the frontier.



- ▶ If goal-state is checked after explored $\rightarrow O(b^{d+1})$

Breadth-First Search

... is not a time and memory friendly choice.

Depth	Nodes	Time	Memory
2	110	.11 ms	107 kB
4	11110	11 ms	10.6 MB
6	10^6	1.1 s	1 GB
8	10^8	2 min	103 GB
10	10^{10}	3 h	10 TB
12	10^{12}	13 days	1 PB
14	10^{14}	3.5 years	99 PB
16	10^{16}	350 years	10 EB

$b = 10$; 1 million nodes/second; 1000 bytes/node



Uniform-Cost Search

- ▶ We need an optimal algorithm for any step cost function, i.e. when the step costs are not equal.
- ▶ Expand the root node, then expand the node with the lowest path cost, ...
- ▶ Implement the frontier as a *priority queue* ordered by path cost function.
- ▶ Add a test to see if a better path is found in a node in the frontier to update priority queue order.

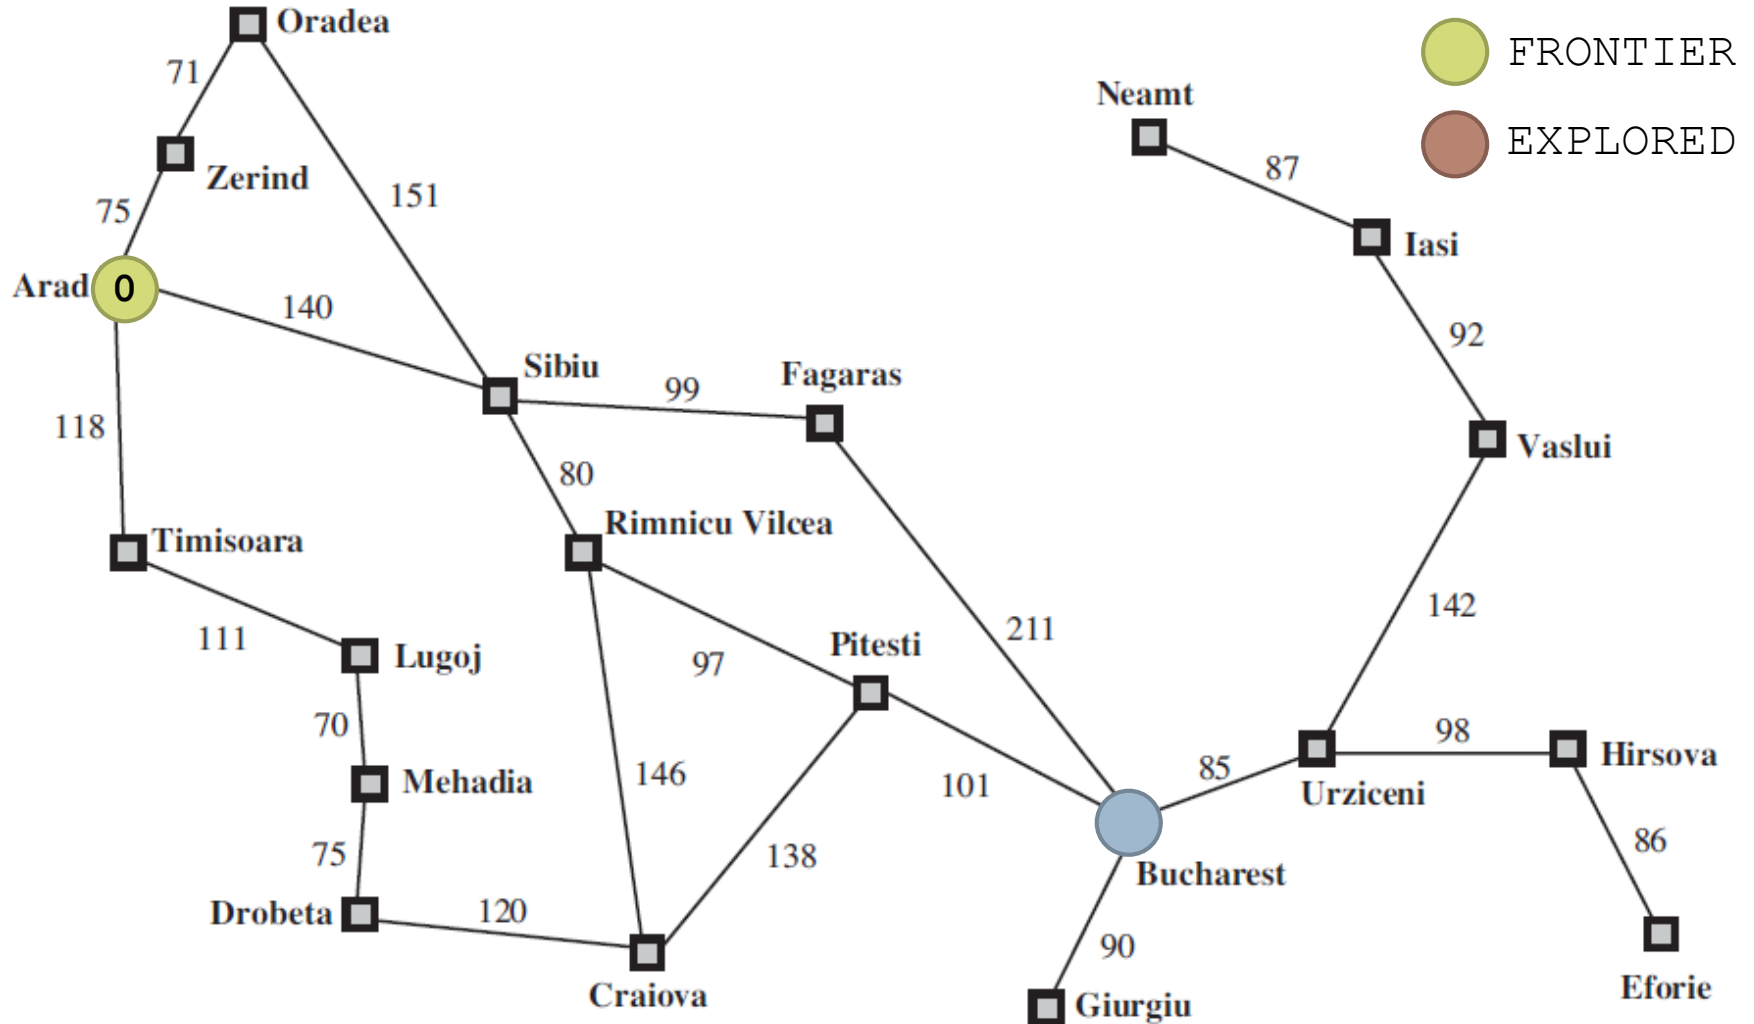


Uniform-Cost Search (for reference)

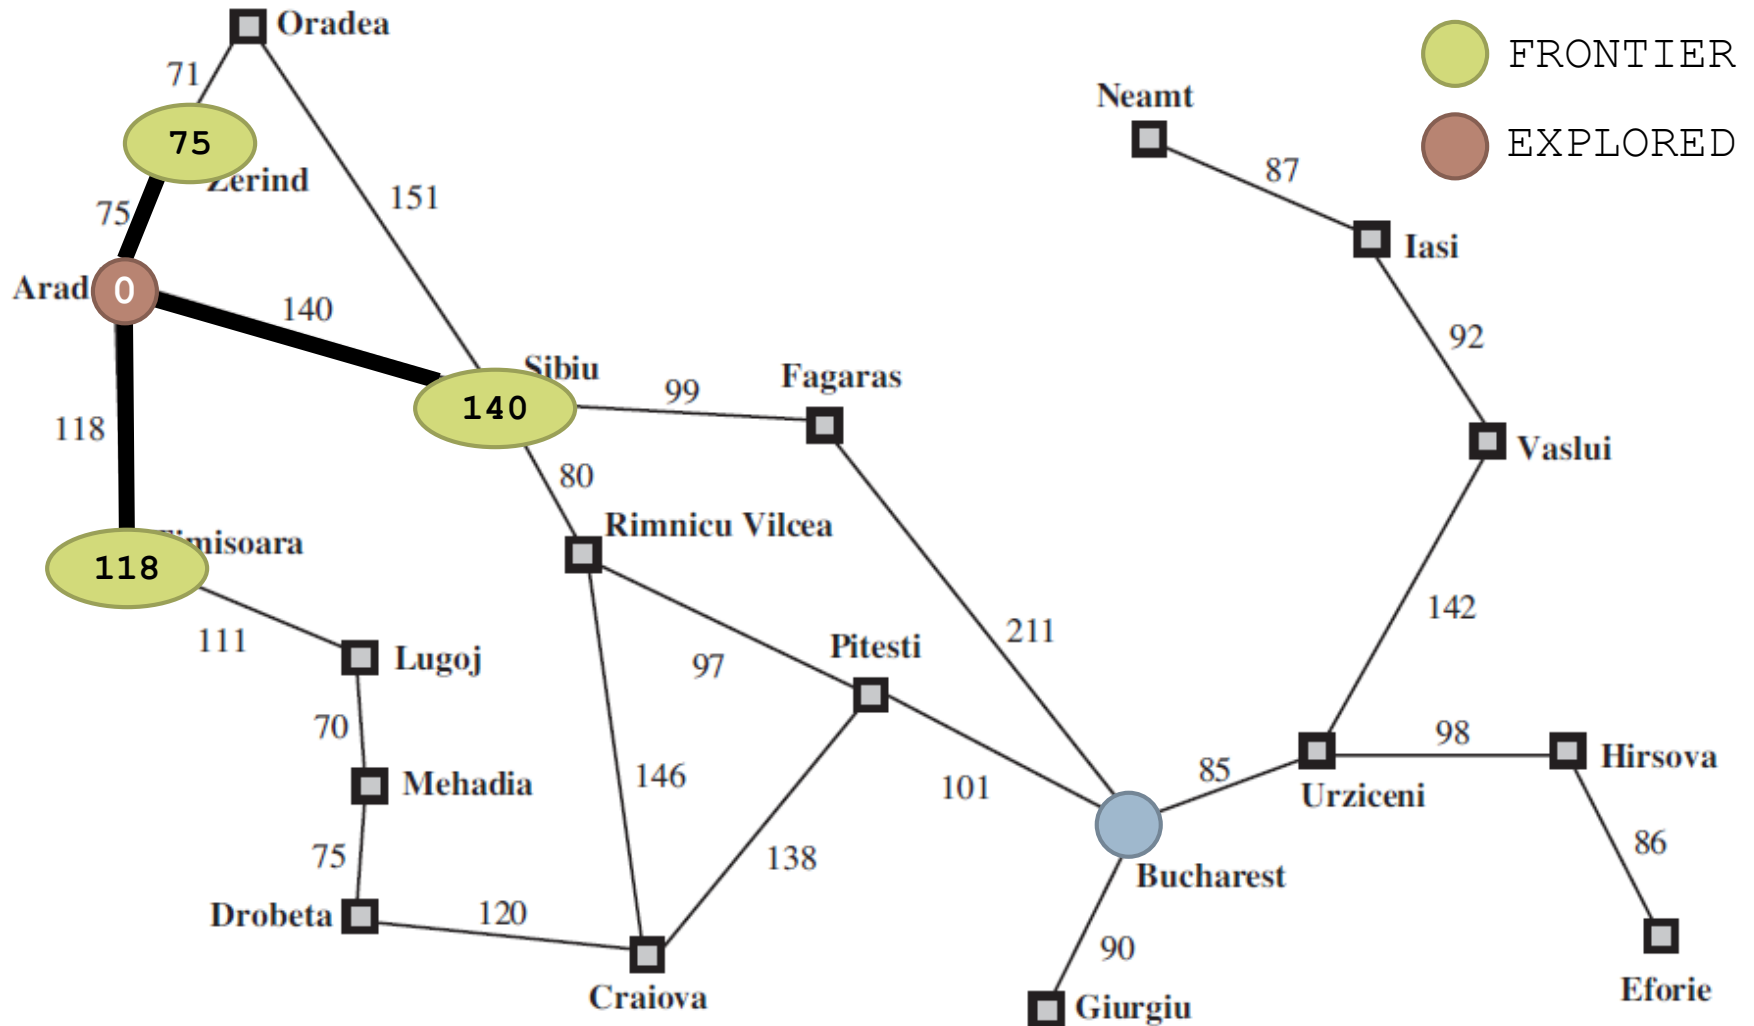
```
function UNIFORM-COST-SEARCH(problem) returns a solution, or failure
    node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
    frontier  $\leftarrow$  a priority queue ordered by PATH-COST, containing only
node
    explored  $\leftarrow$  an empty set
    loop do
        if EMPTY?( frontier) then return failure
        node  $\leftarrow$  POP( frontier ) /* chooses the lowest-cost node in
frontier */
        if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
        add node.STATE to explored
        for each action in problem.ACTIONS(node.STATE) do
            child  $\leftarrow$  CHILD-NODE(problem, node, action)
            if child.STATE is not in explored or frontier then
                frontier  $\leftarrow$  INSERT(child, frontier)
            else if child.STATE is in frontier with higher PATH-COST then
                replace that frontier node with child
```



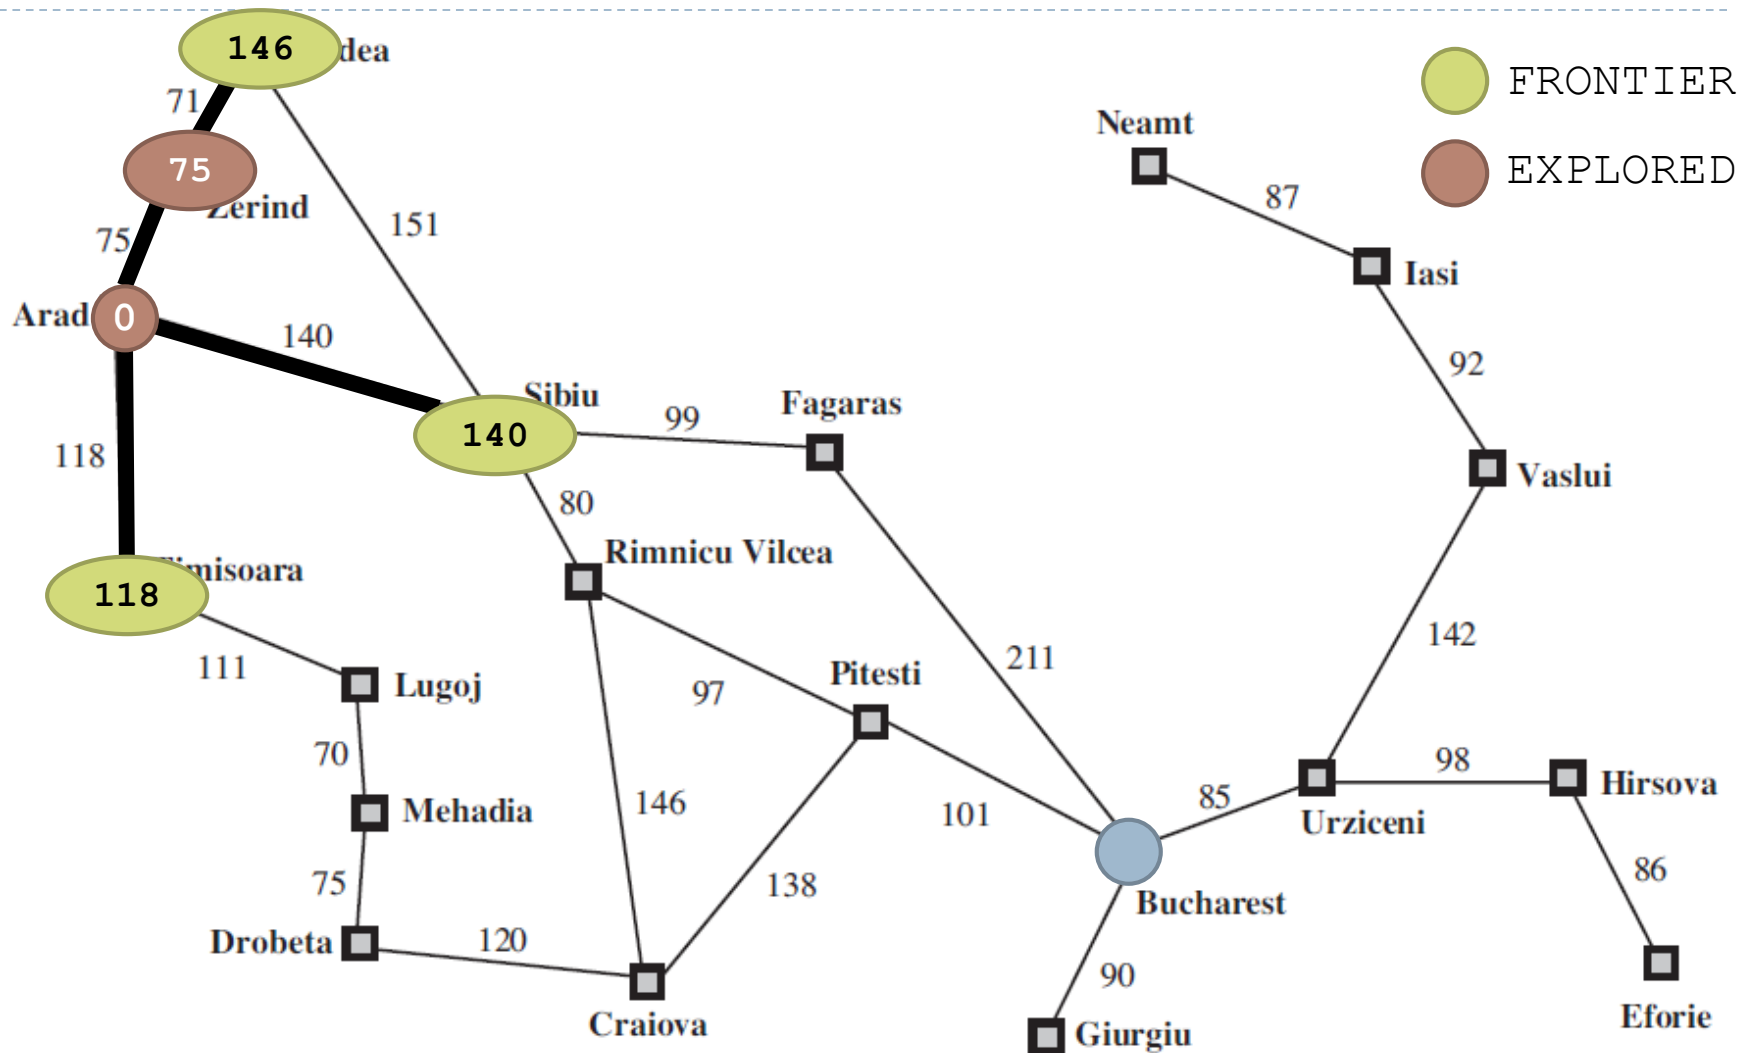
Uniform-Cost Search for Route Finding



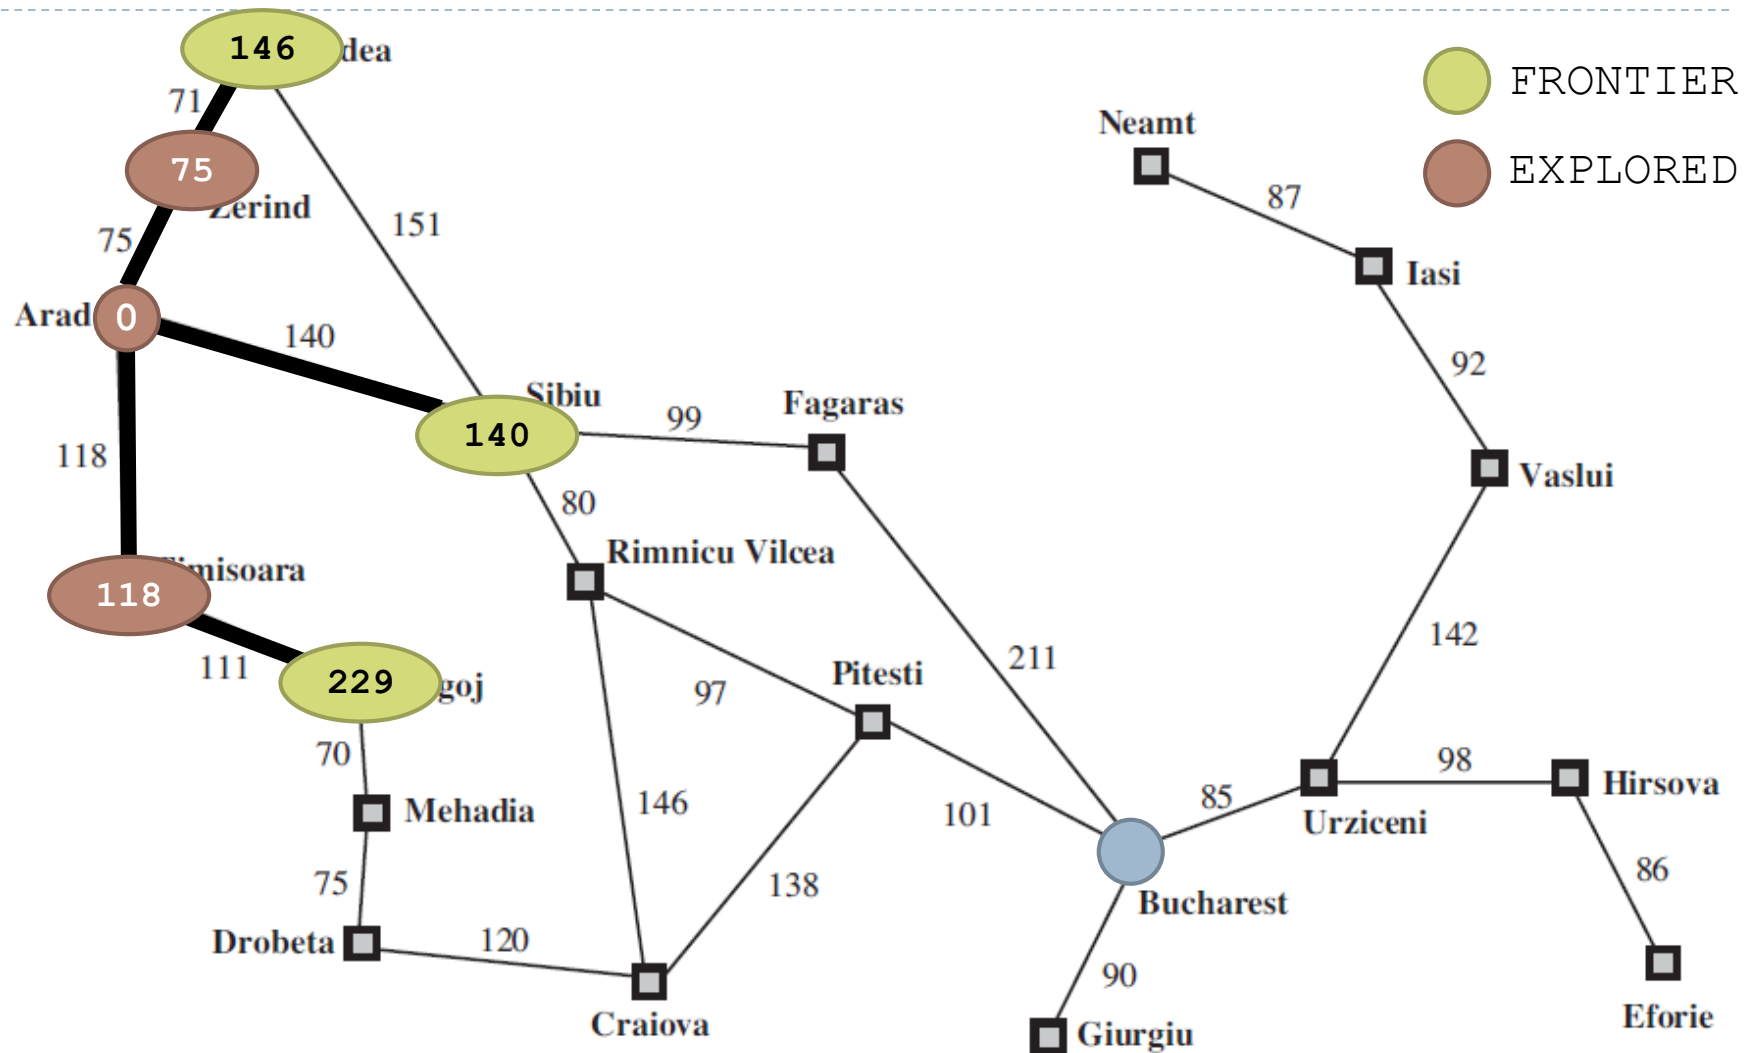
Uniform-Cost Search for Route Finding



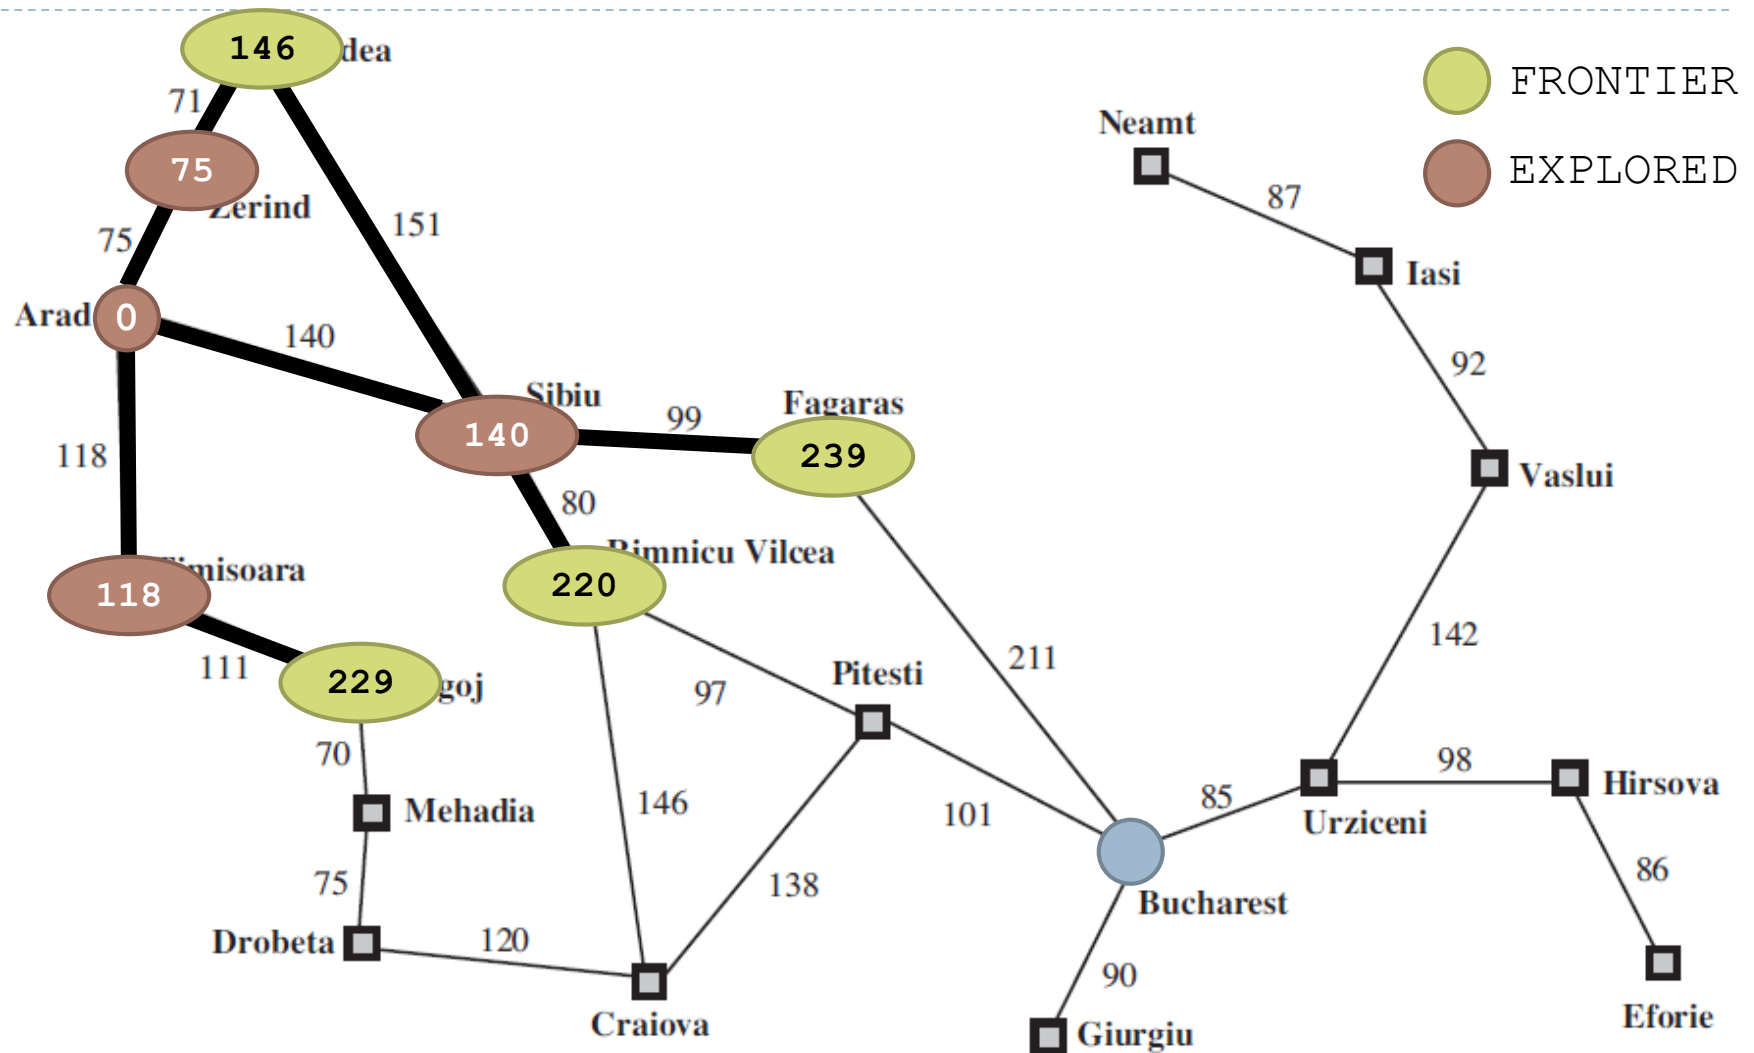
Uniform-Cost Search for Route Finding



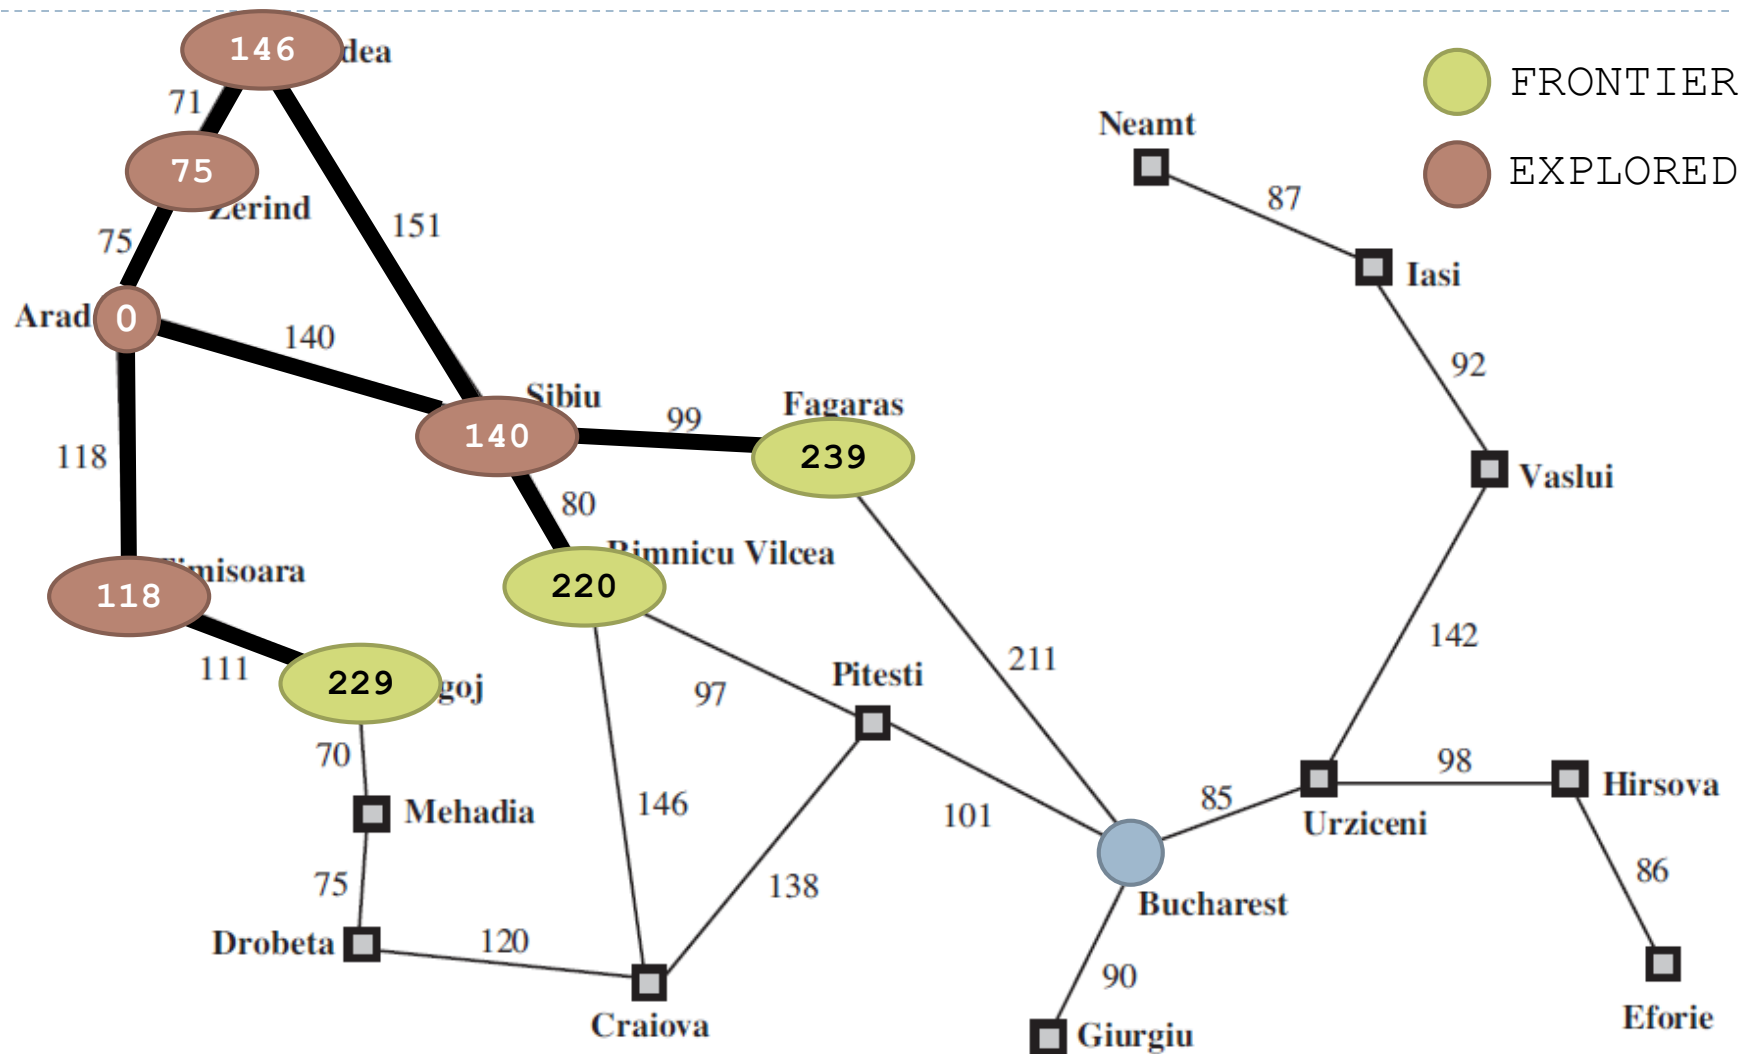
Uniform-Cost Search for Route Finding



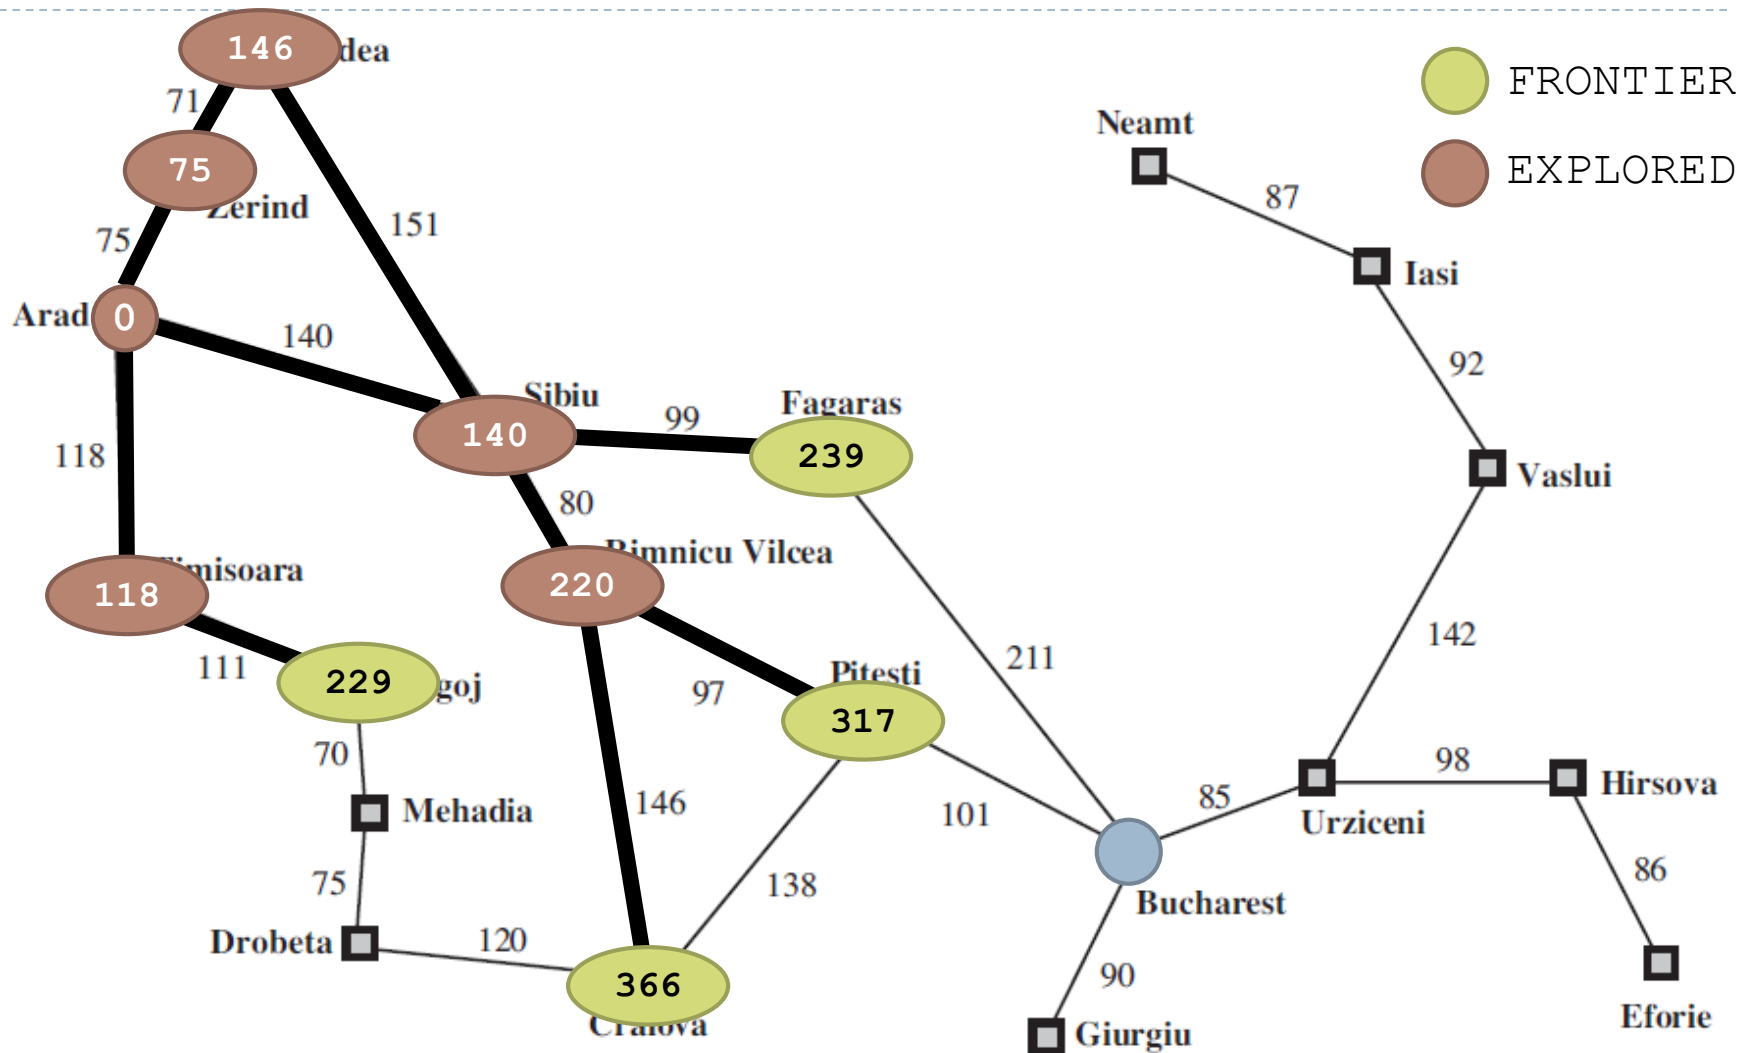
Uniform-Cost Search for Route Finding



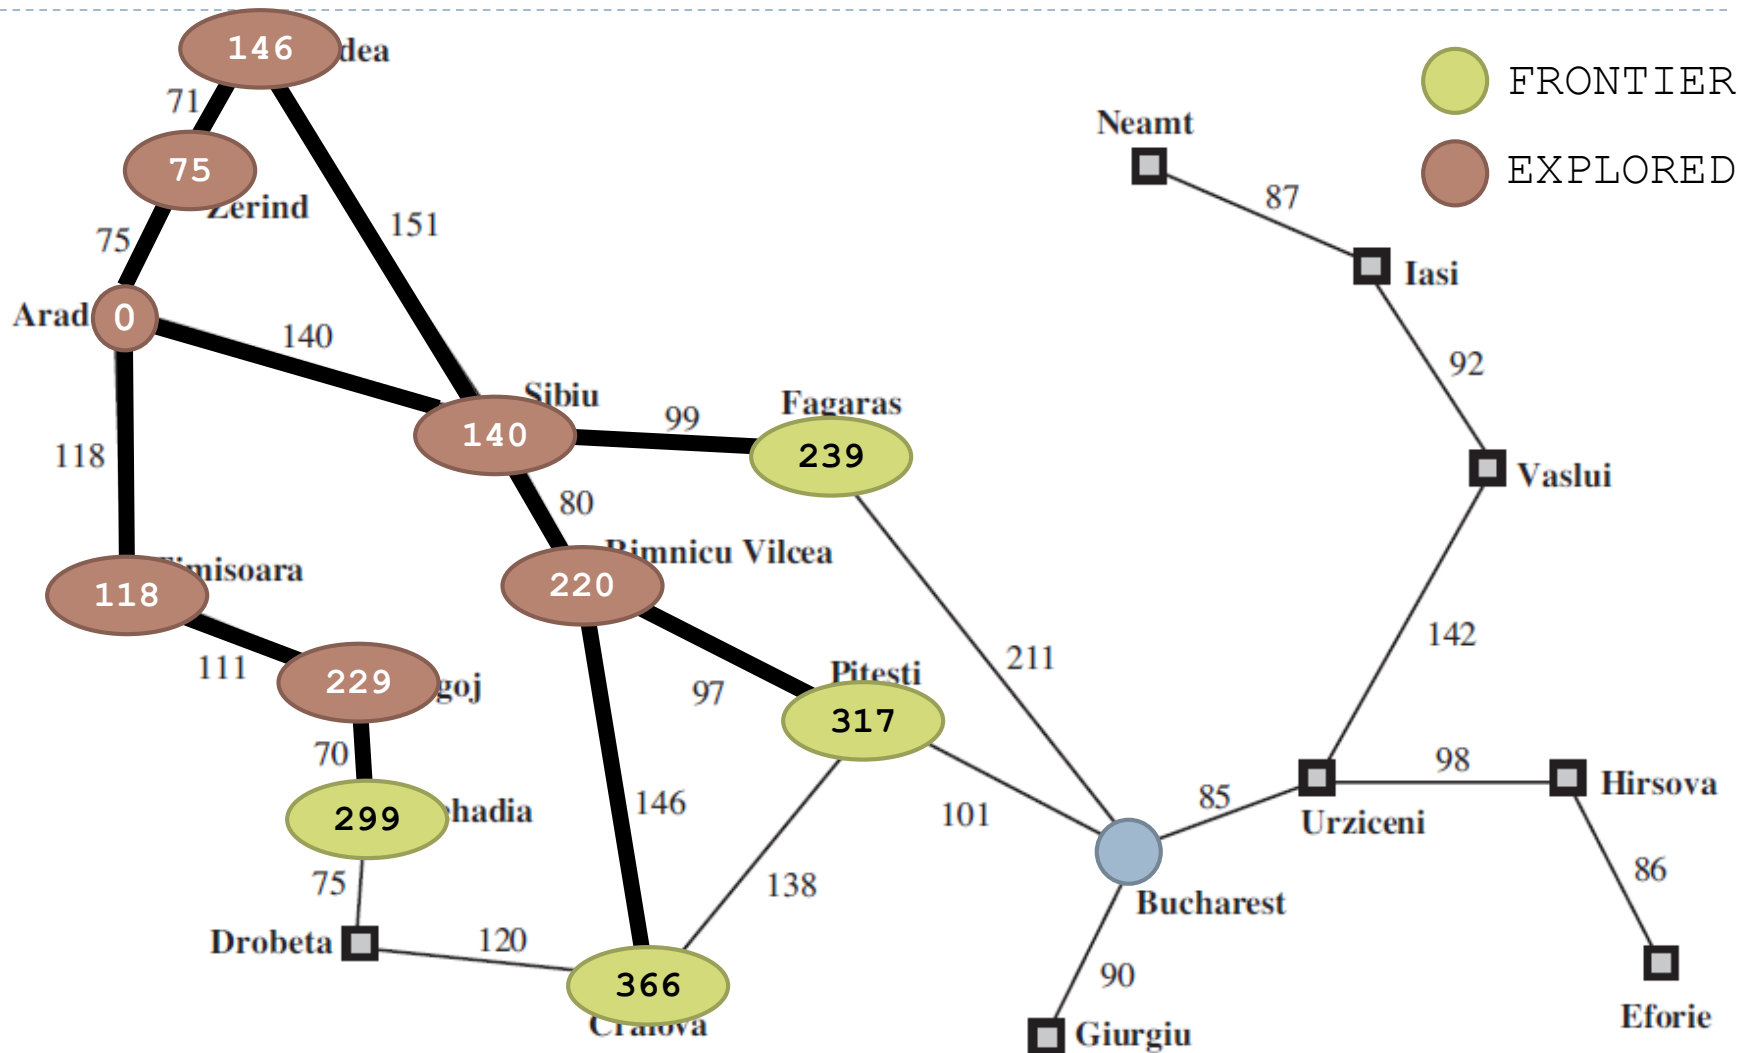
Uniform-Cost Search for Route Finding



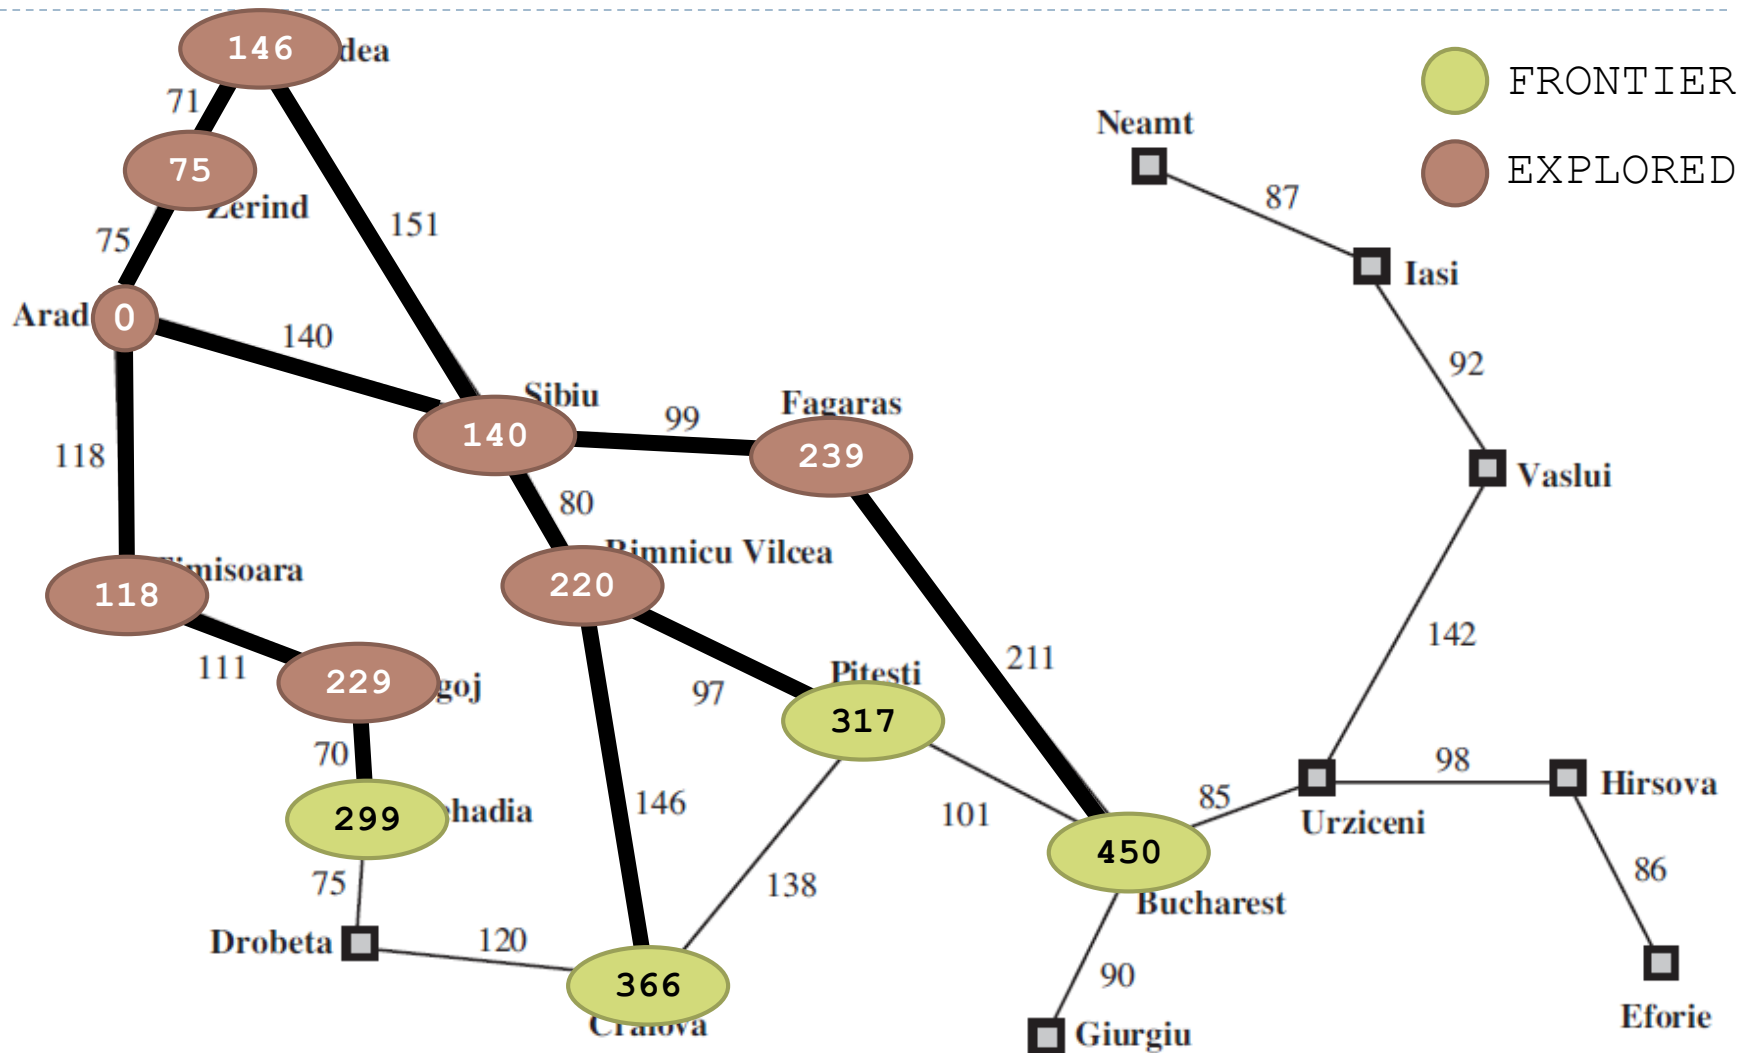
Uniform-Cost Search for Route Finding



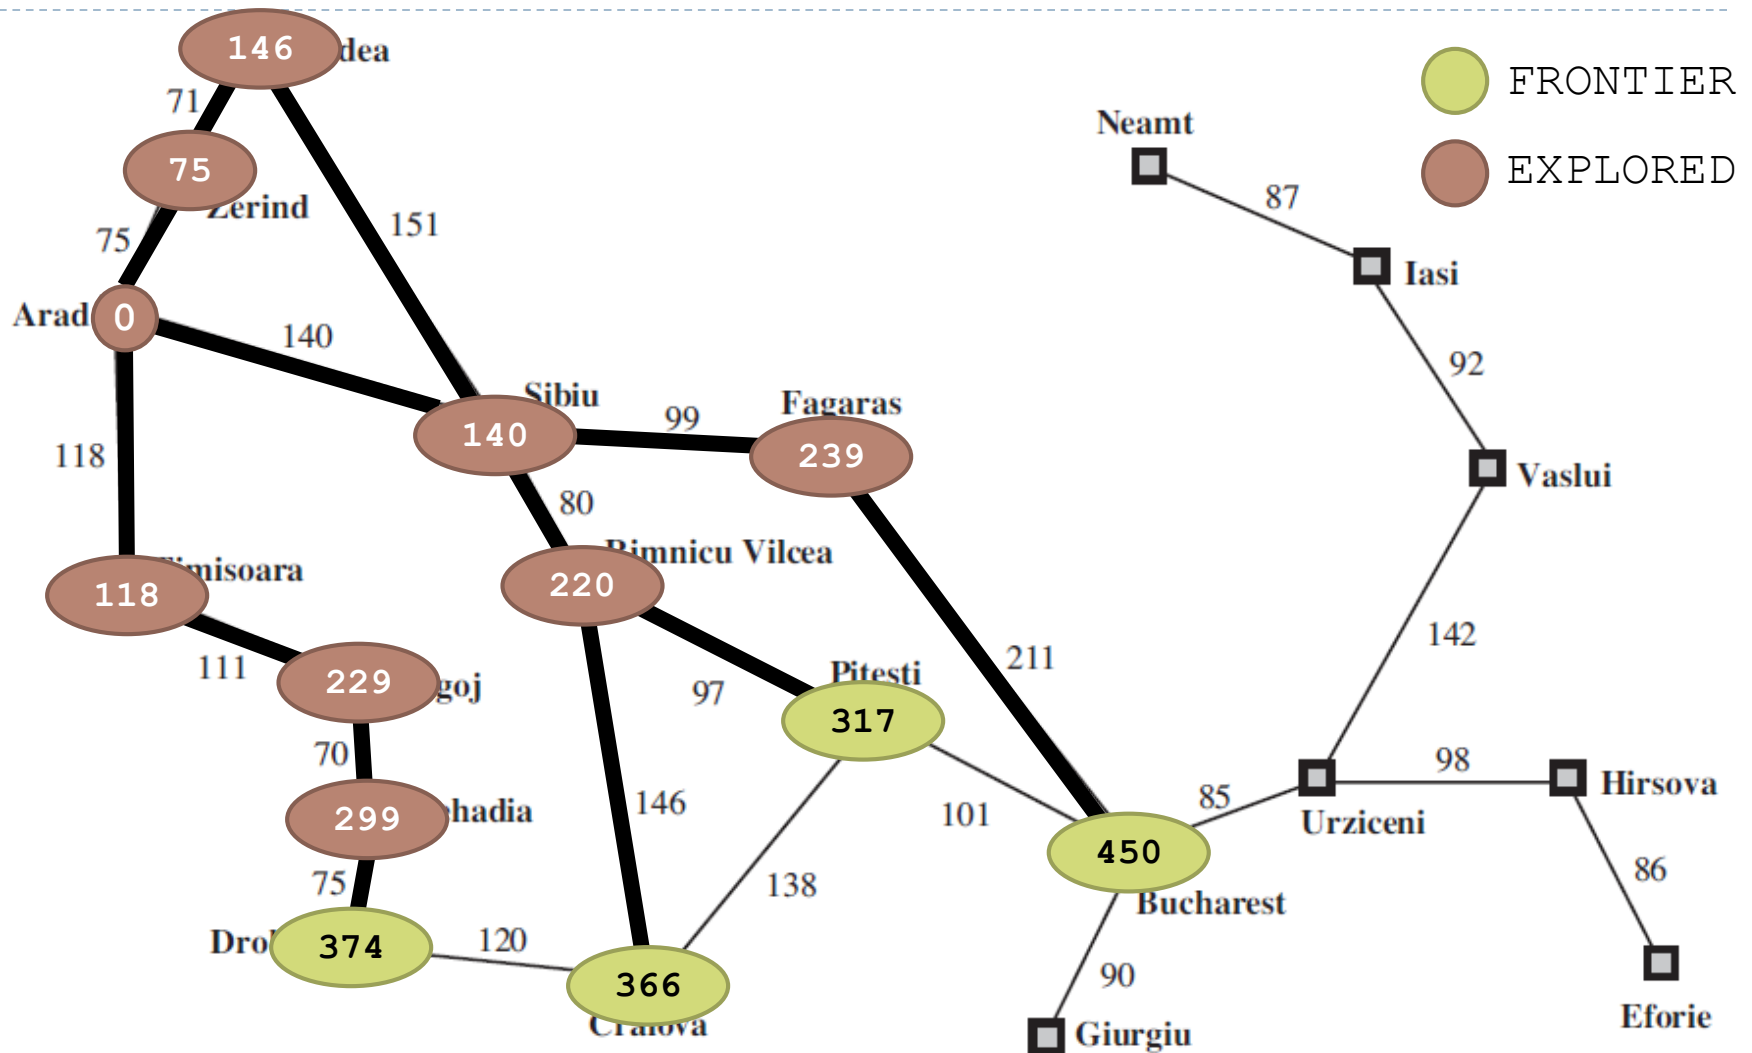
Uniform-Cost Search for Route Finding



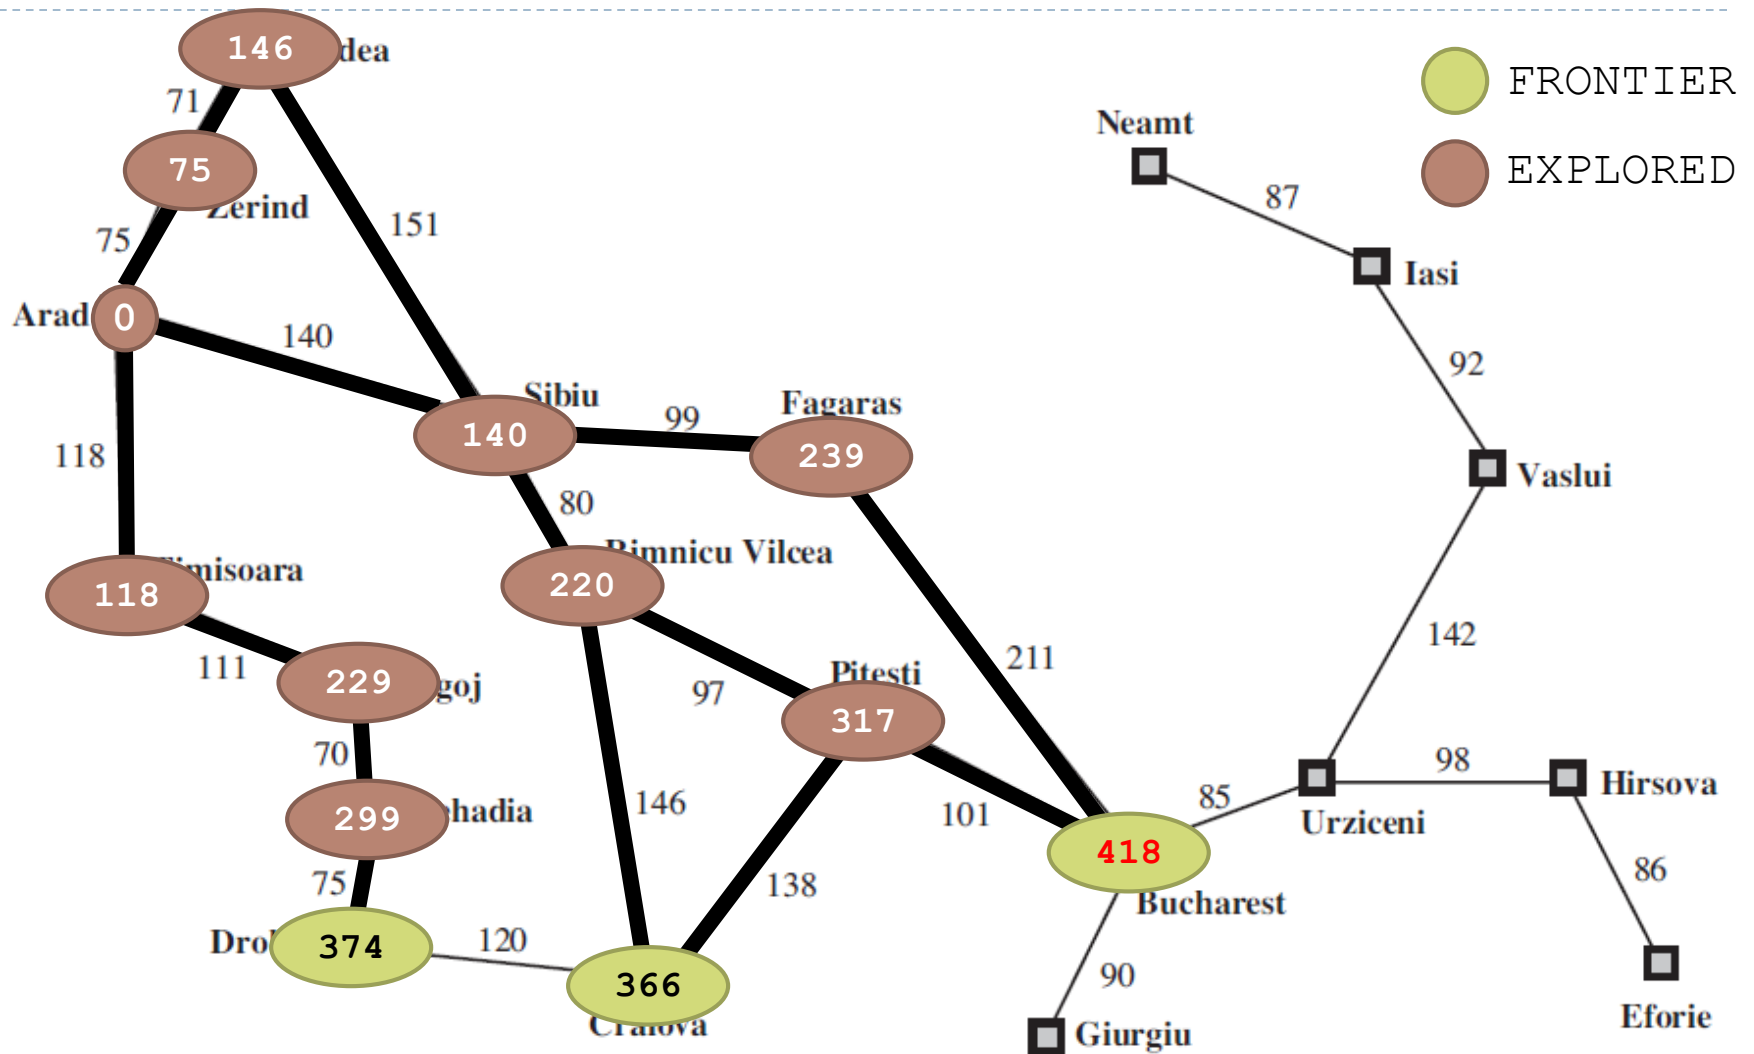
Uniform-Cost Search for Route Finding



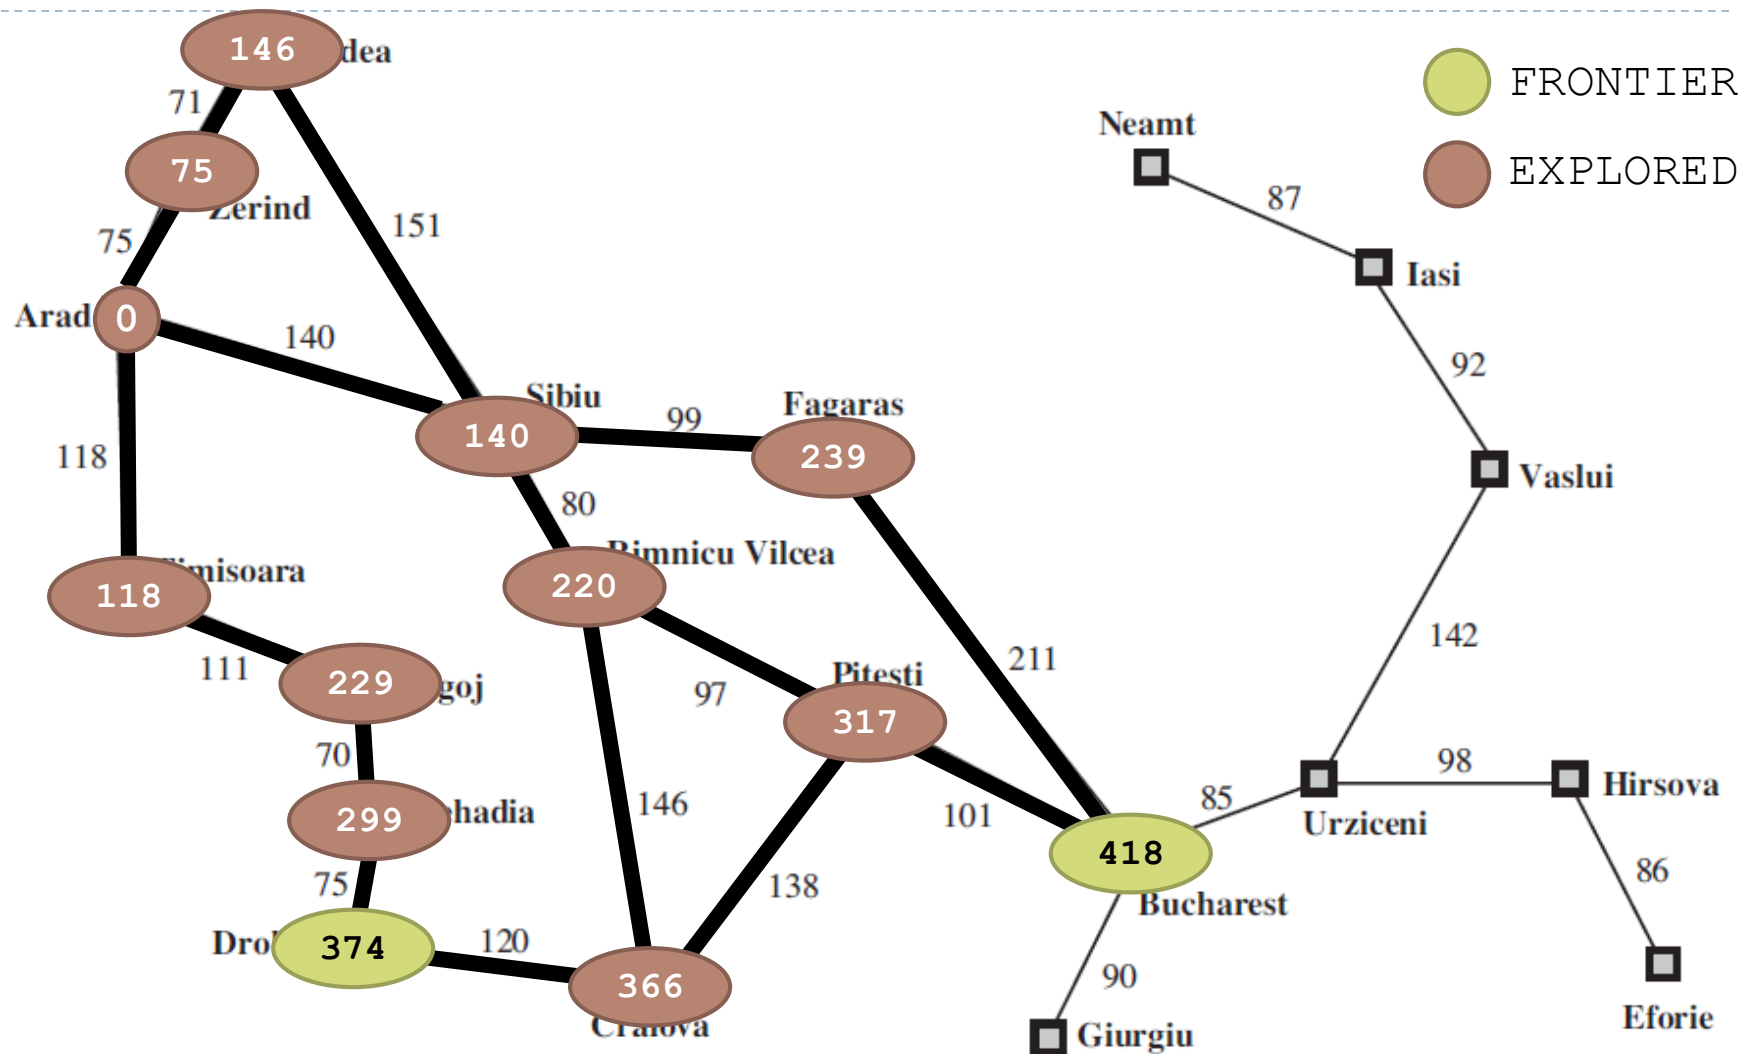
Uniform-Cost Search for Route Finding



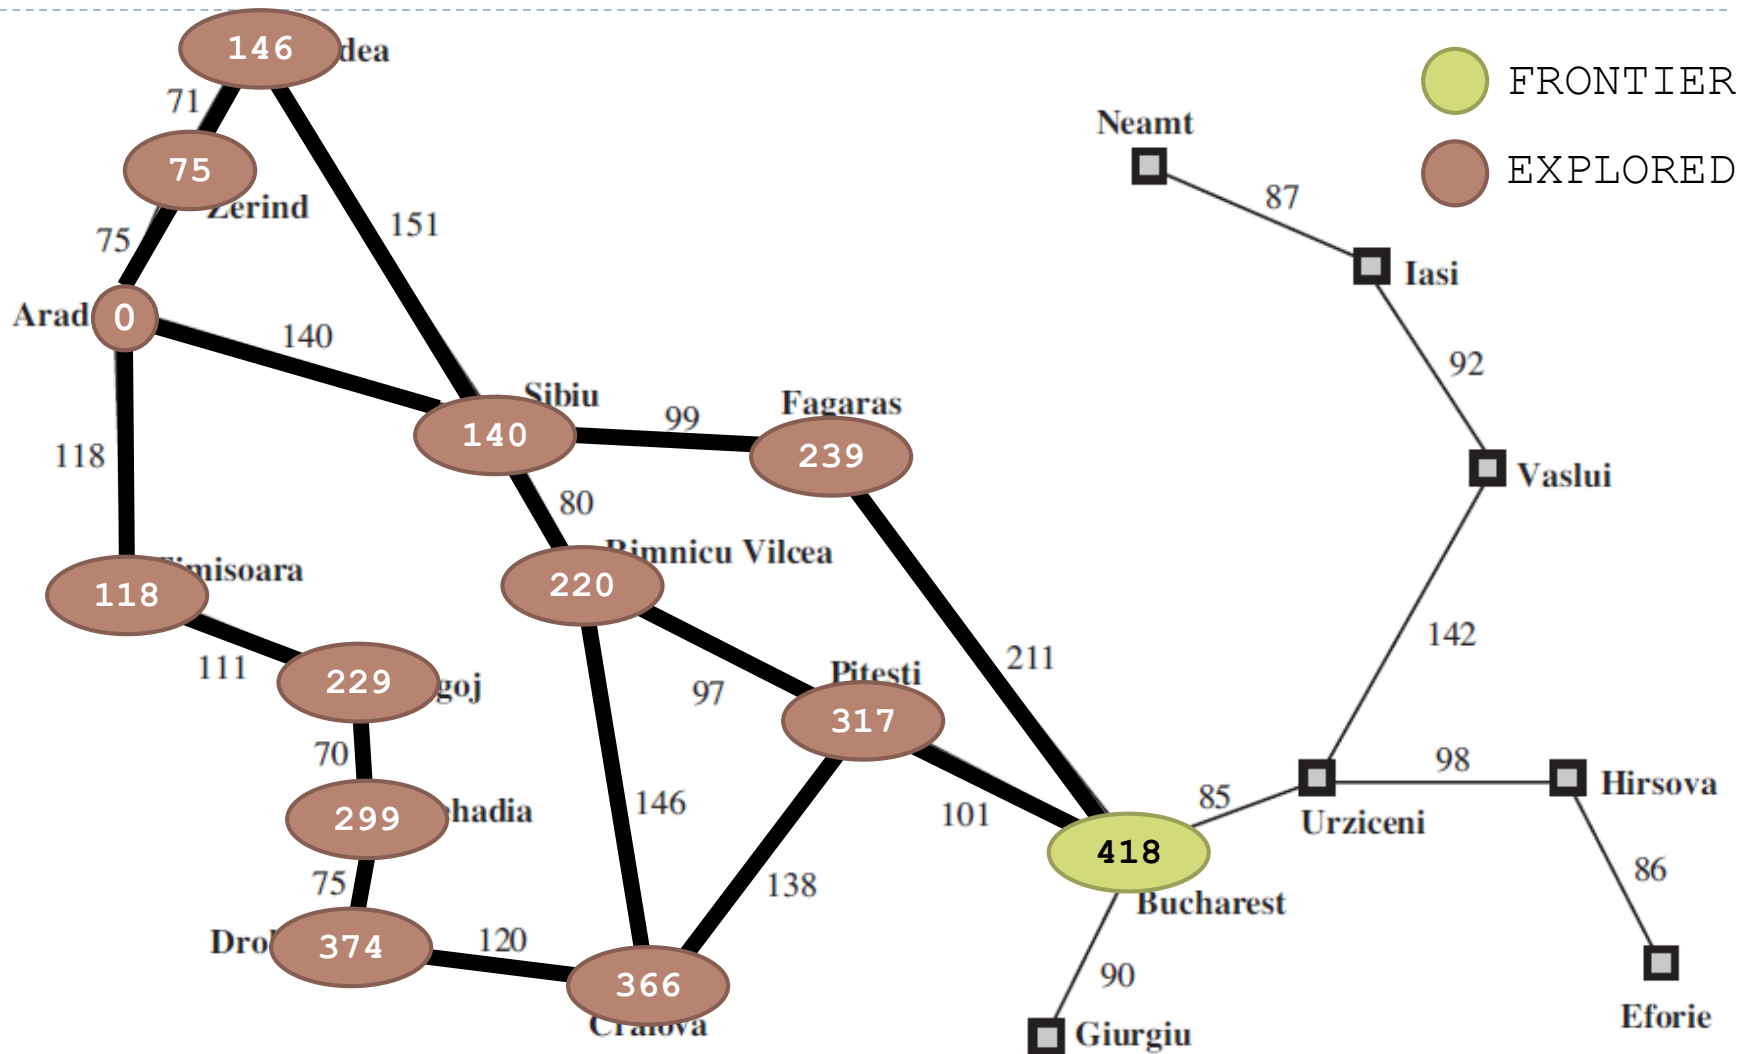
Uniform-Cost Search for Route Finding



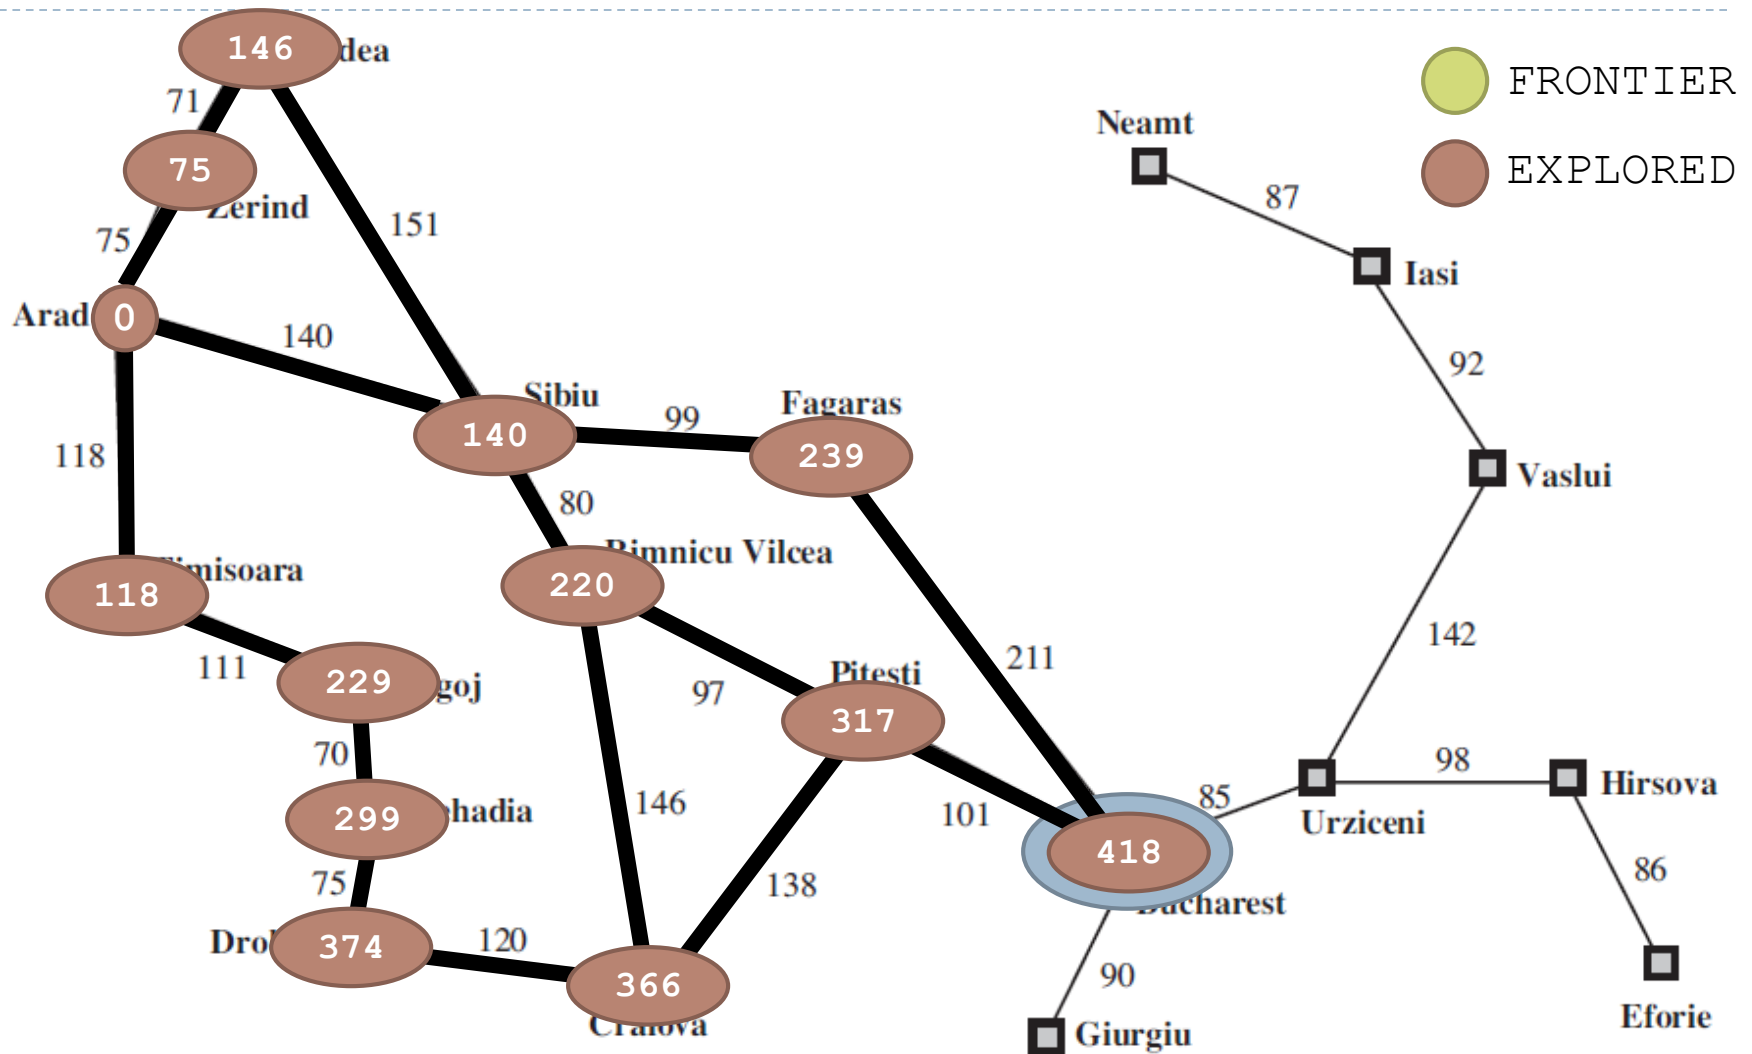
Uniform-Cost Search for Route Finding



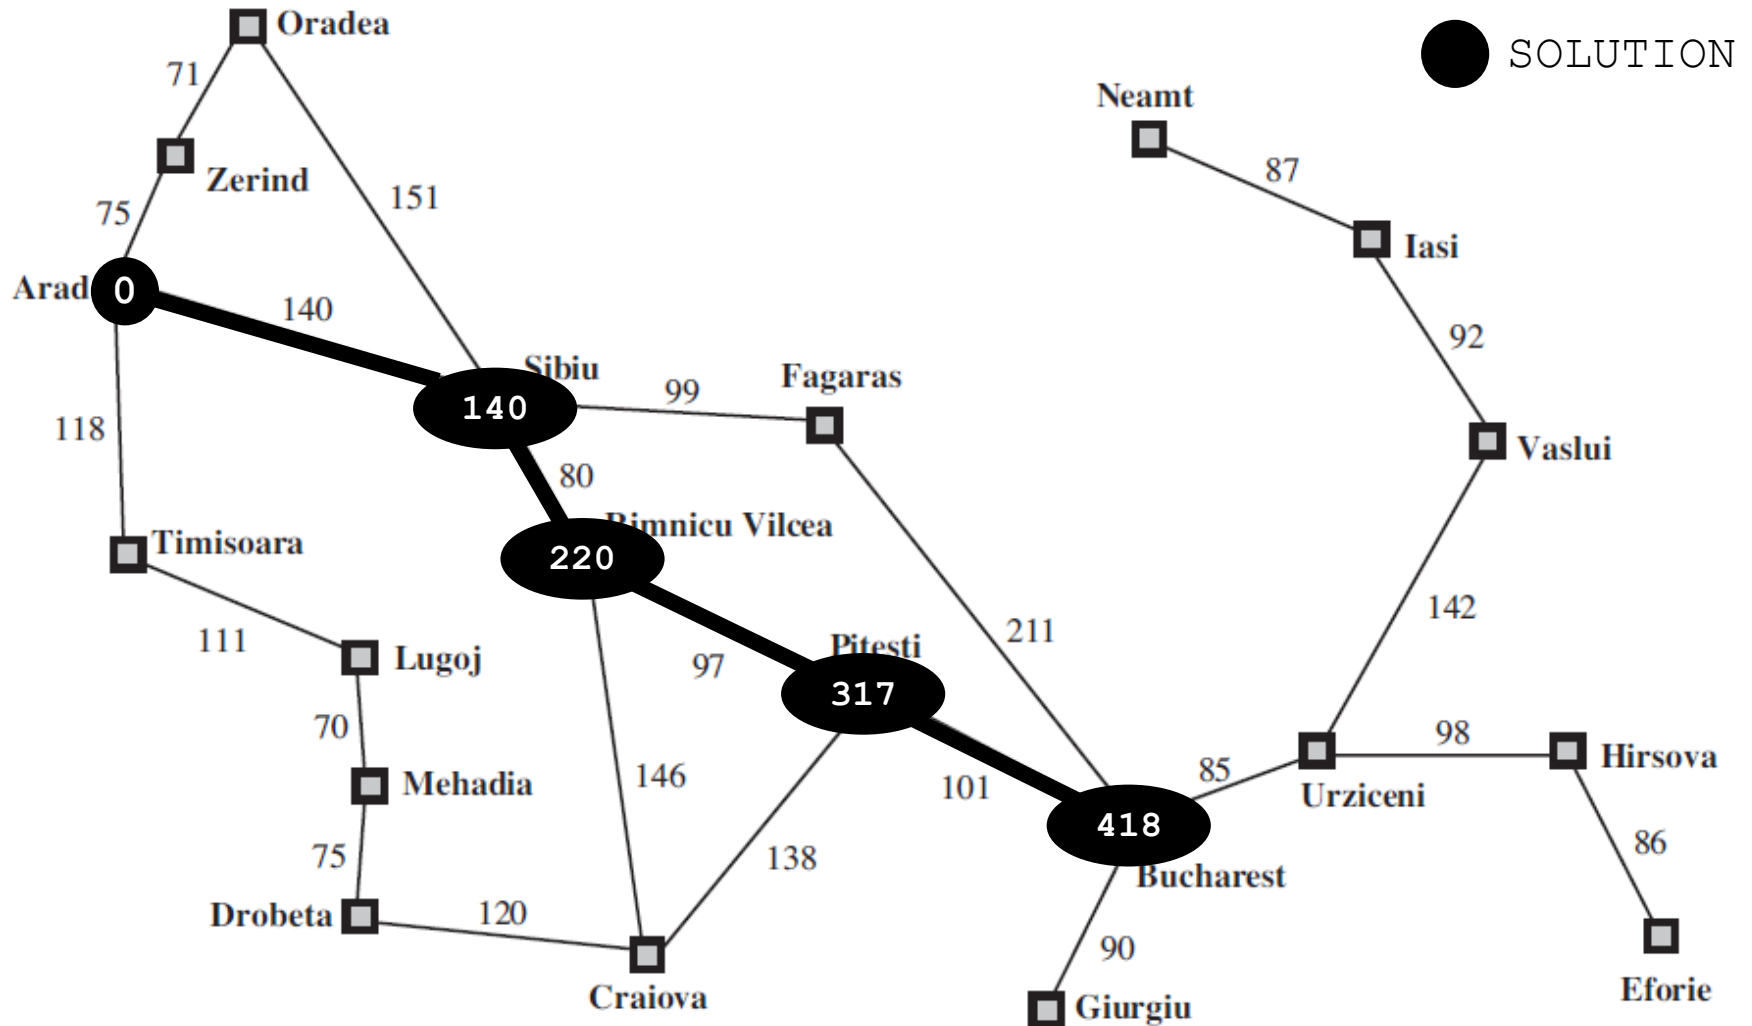
Uniform-Cost Search for Route Finding



Uniform-Cost Search for Route Finding



Uniform-Cost Search for Route Finding



Uniform-Cost Search


- ▶ It is optimal, i.e. the first solution we find is the minimum cost one.
- ▶ It is complete if every step cost exceeds some small positive constant, ϵ

There should not be zero or negative cost loops

- ▶ Its space and time complexity is

$O(b^{l + \lceil C^* / \epsilon \rceil})$, where C^* is the optimal path cost and ϵ is the minimal action cost.

In case the step costs are all equal, this is equivalent to $O(b^{l+d})$

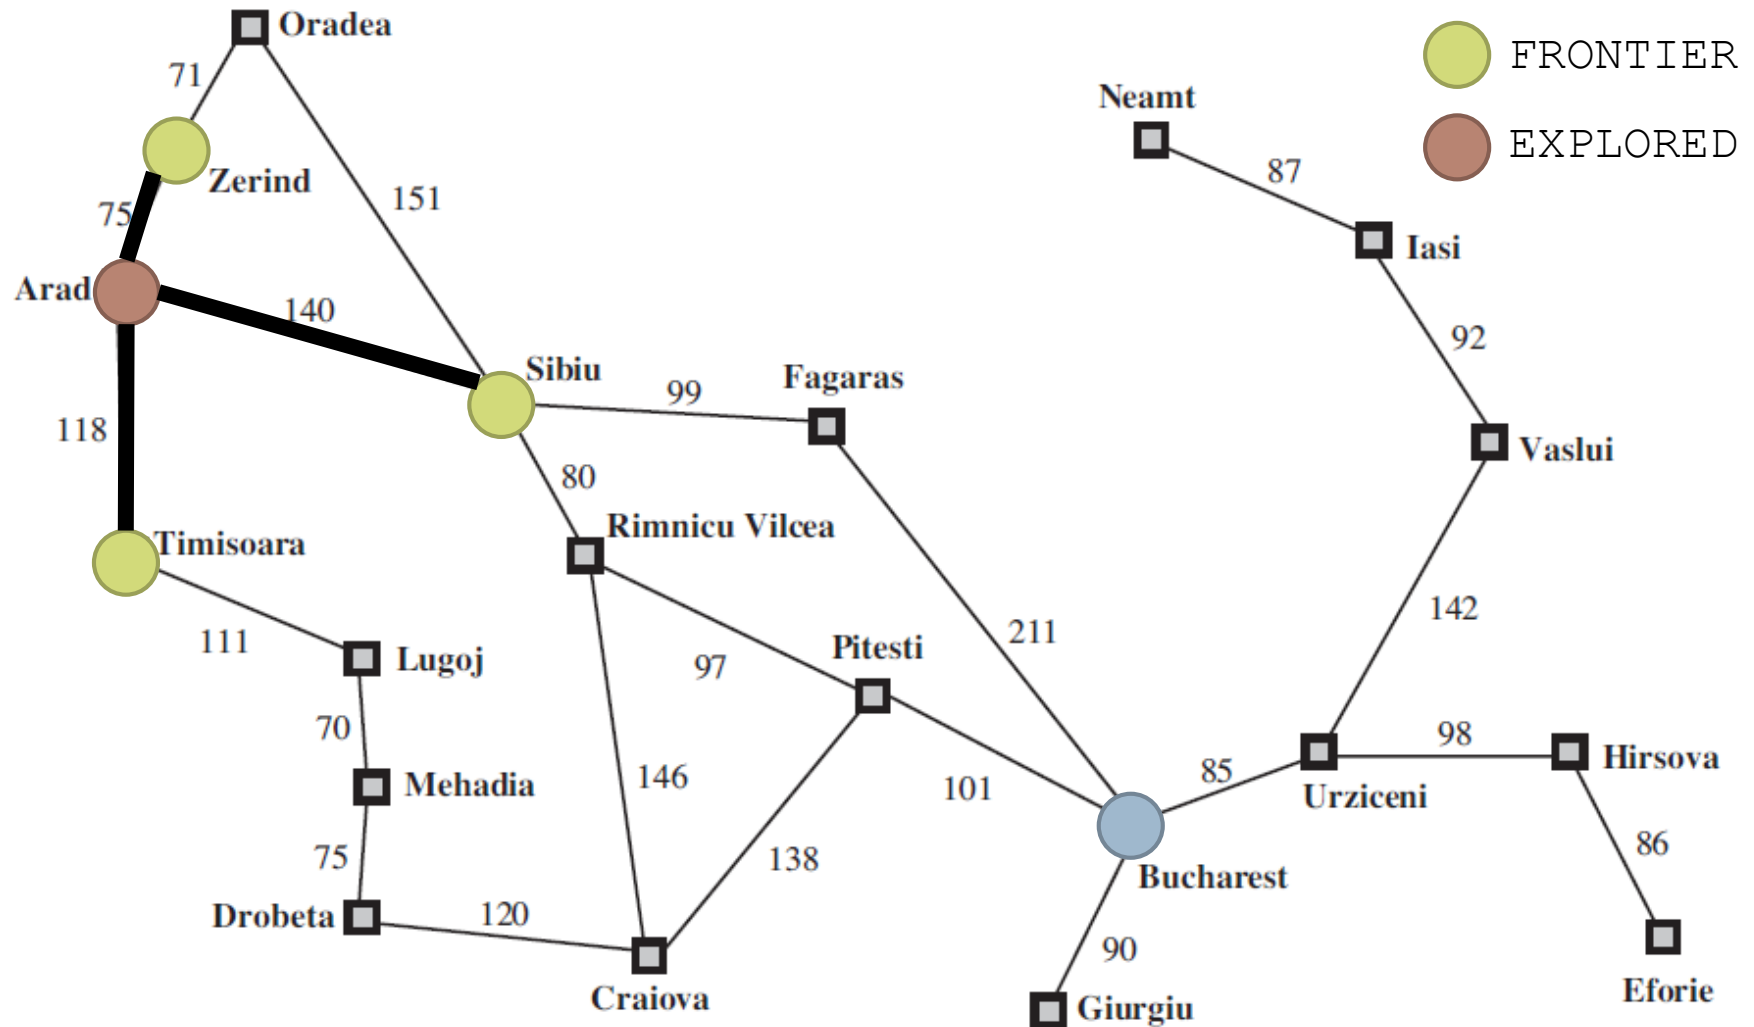


Depth-First Search

- ▶ Expand the root node, then keep expanding the deepest unexplored node.
- ▶ Implement the frontier as a LIFO queue (a stack).

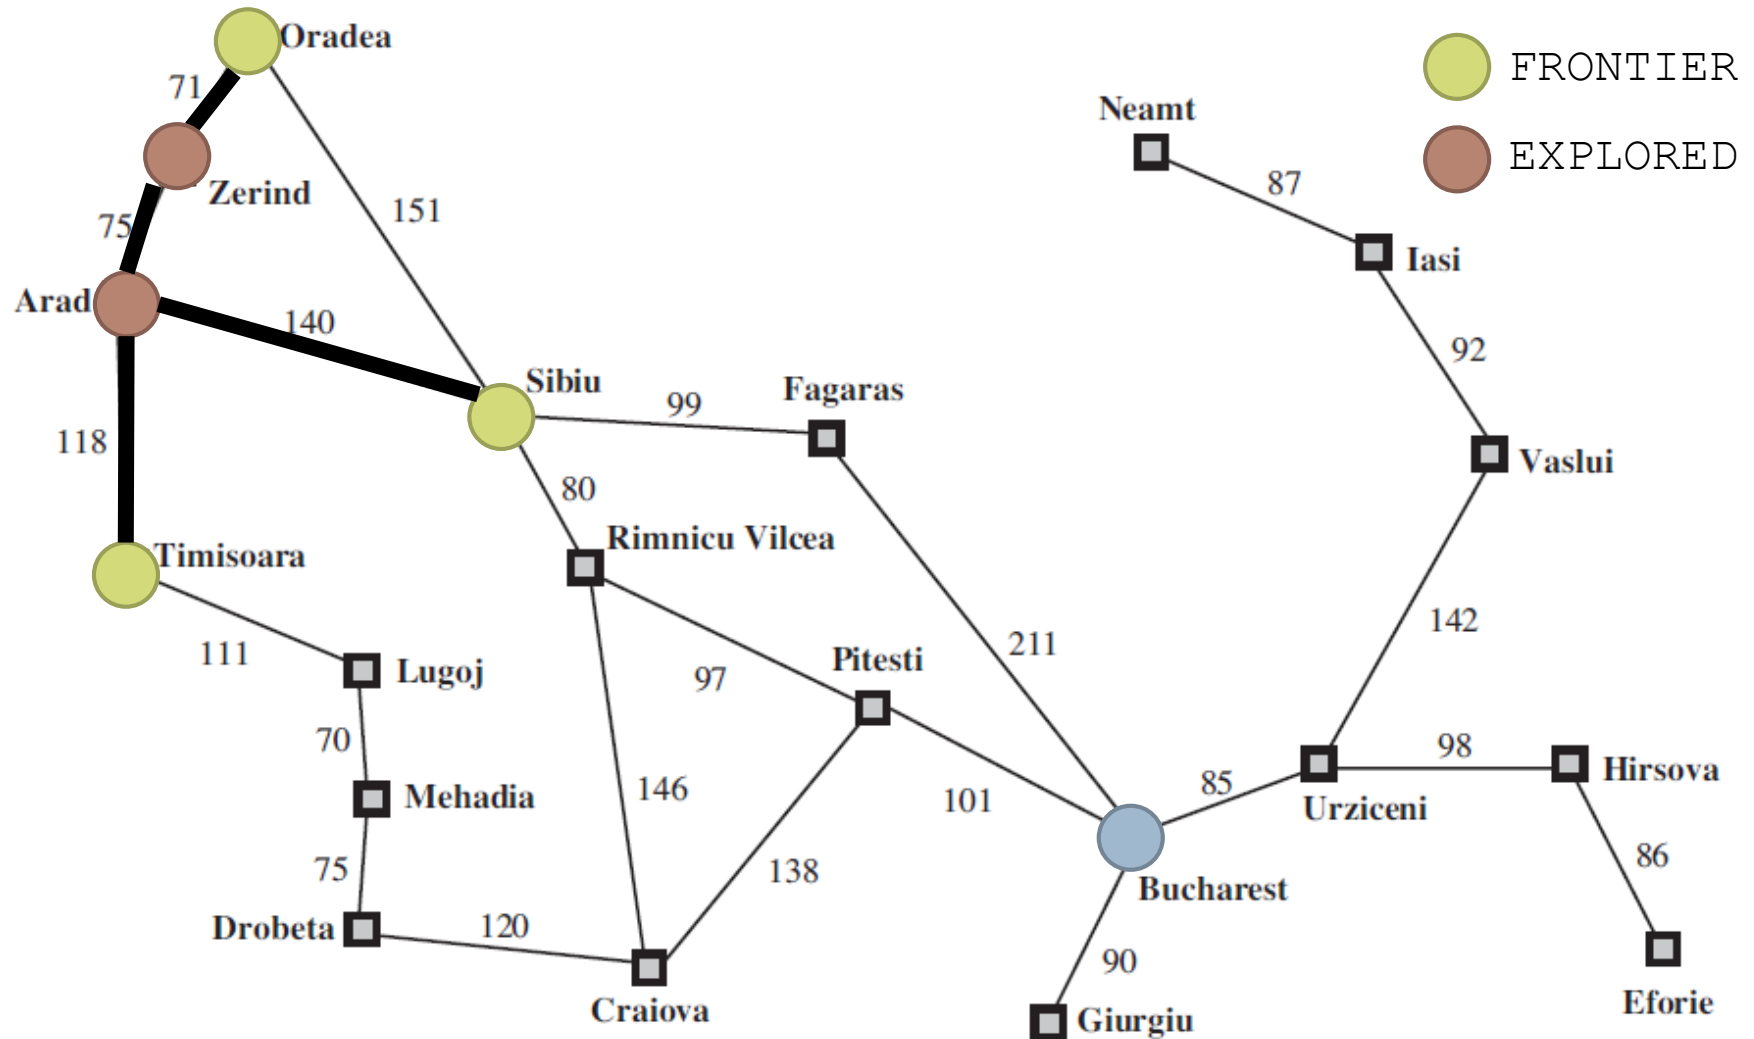


Depth-First Search for Route Finding




► Break the ties by expanding the node starting with alphabetically last character.

Depth-First Search for Route Finding




► Break the ties by expanding the node starting with alphabetically last character.

Depth-First Search for Route Finding



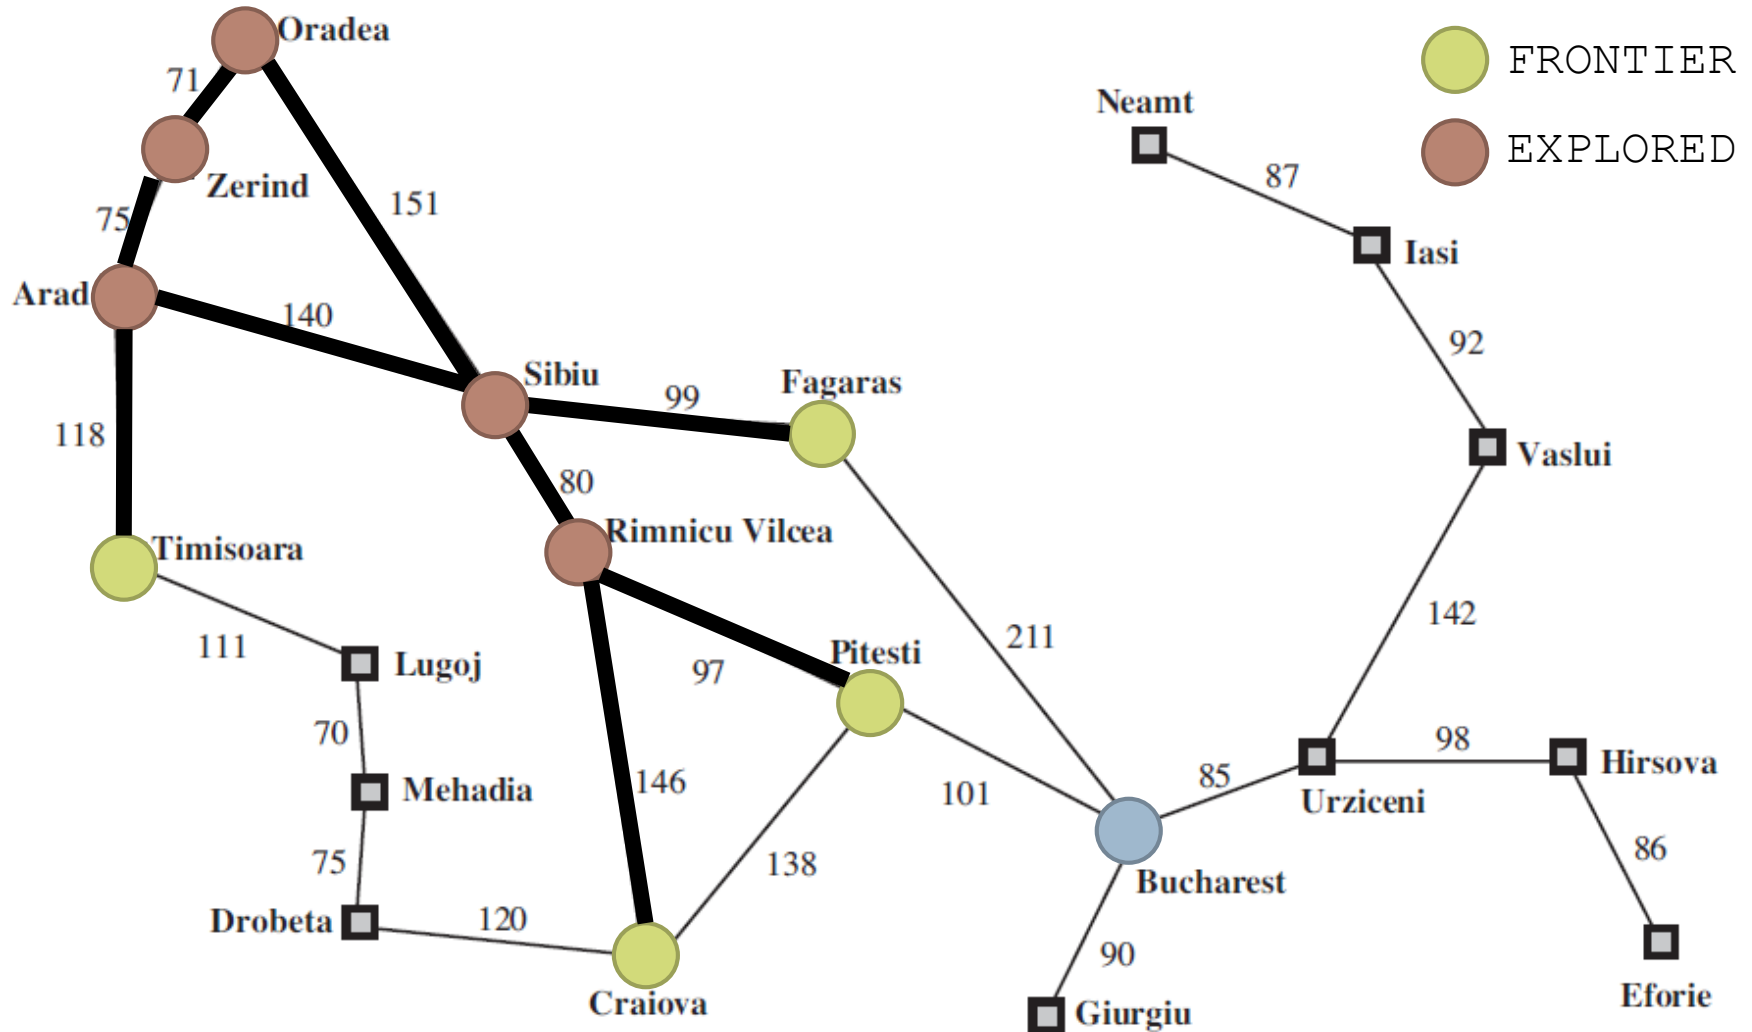
► Break the ties by expanding the node starting with alphabetically last character.

Depth-First Search for Route Finding



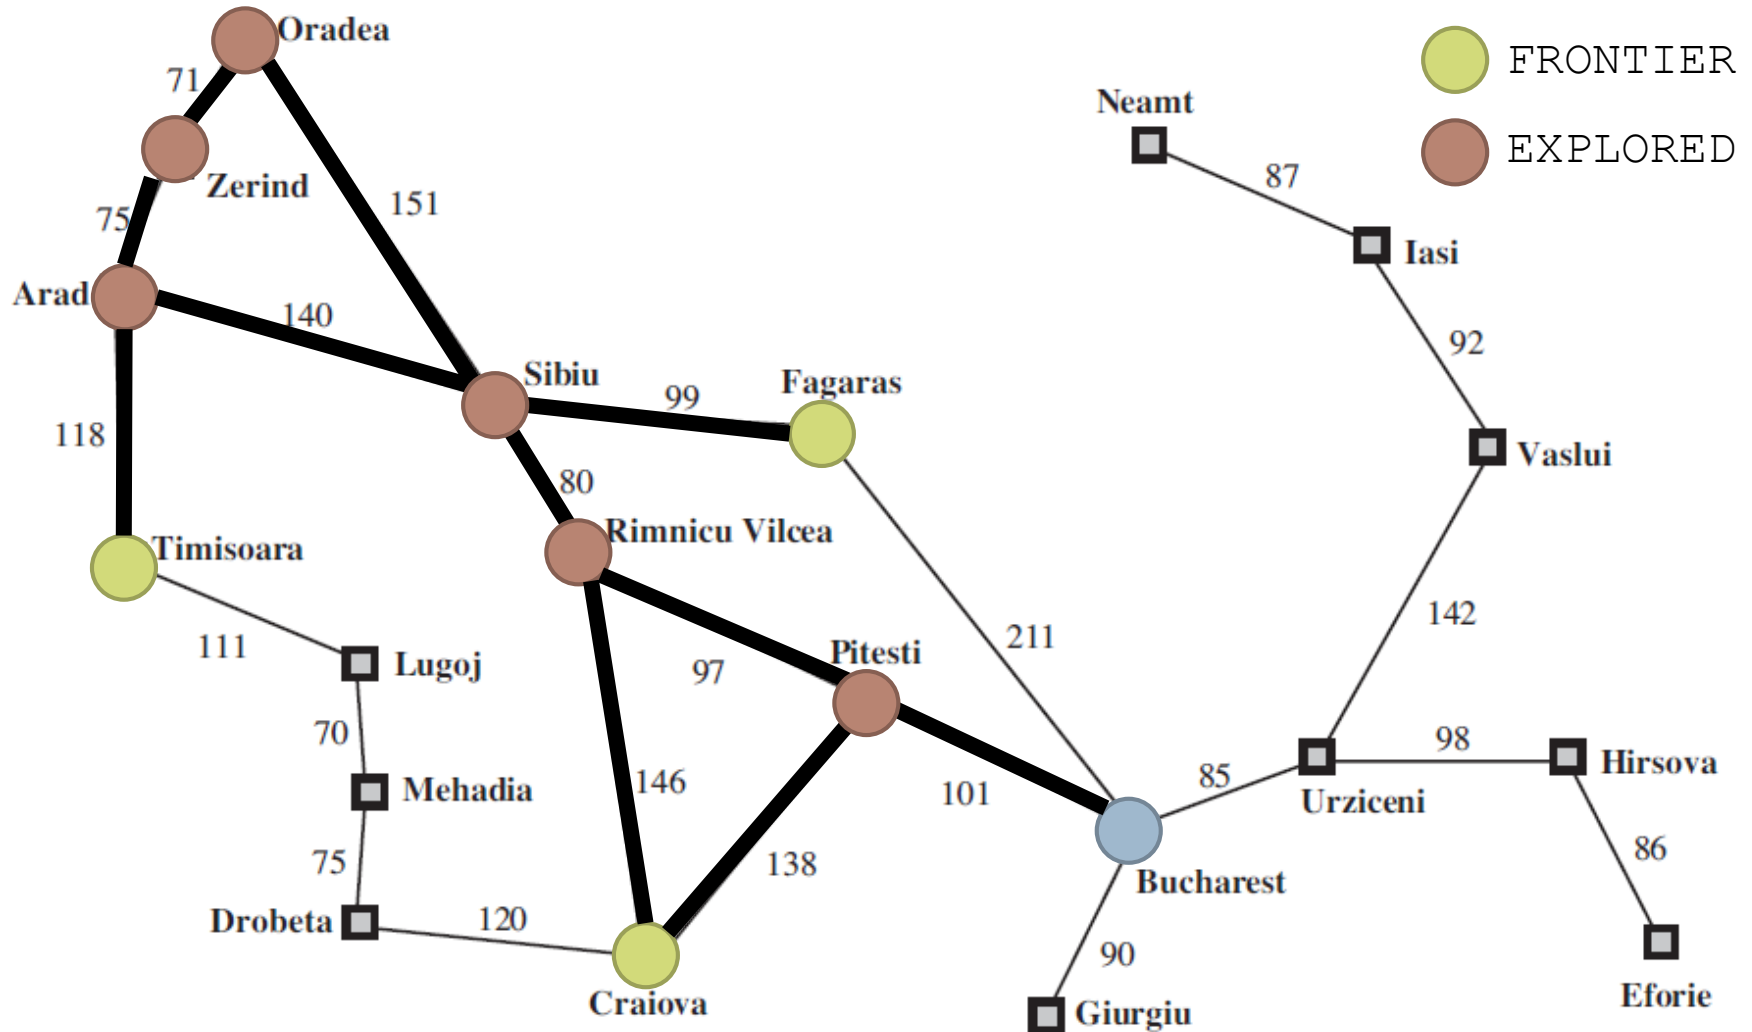
► Break the ties by expanding the node starting with alphabetically last character.

Depth-First Search for Route Finding




► Break the ties by expanding the node starting with alphabetically last character.

Depth-First Search for Route Finding



► Break the ties by expanding the node starting with alphabetically last character.


Depth-First Search for Route Finding



► Is it the minimum cost path?


Depth-First Search

- ▶ It is not optimal.
- ▶ It is not complete for an infinite search tree. It never finds the solution.
- ▶ Its time complexity is $O(b^m)$,
b is the branching factor, m is the maximum depth of the tree.
- ▶ Its space complexity is $O(bm)$,
and only m nodes in the frontier.
- ▶ Its low space complexity makes it a good candidate for actual implementations.




Search comparison


Breadth-first
search



Uniform-cost
search



Depth-first
search




For each search method, tell in what order the tree is expanded?
For each search method, tell if it is optimal?
(solution can be at any node)




Search comparison


Breadth-first
search



Uniform-cost
search



Depth-first
search



For each search method, tell in what order the tree is expanded?
Let's assume blue circles are solutions, then depth-first search finds a higher cost solution before a minimum cost one.



Variants of the Above

- ▶ **Depth-Limited Search**

- ▶ Limit the depth to a certain level to avoid the failure of depth-first search in infinite state spaces.

- ▶ **Iterative Deepening Depth-First Search**

- ▶ Iteratively increase the depth level limit.


- ▶ **Bidirectional Search**

- ▶ Start from the initial and goal states and expand both, check for intersection of frontiers.



Home-Exercise

- ▶ Apply Breadth-first, Depth-first, Uniform-cost.



- ▶ Break the ties by expanding the node starting with alphabetically first character.